

# Physical Database Design



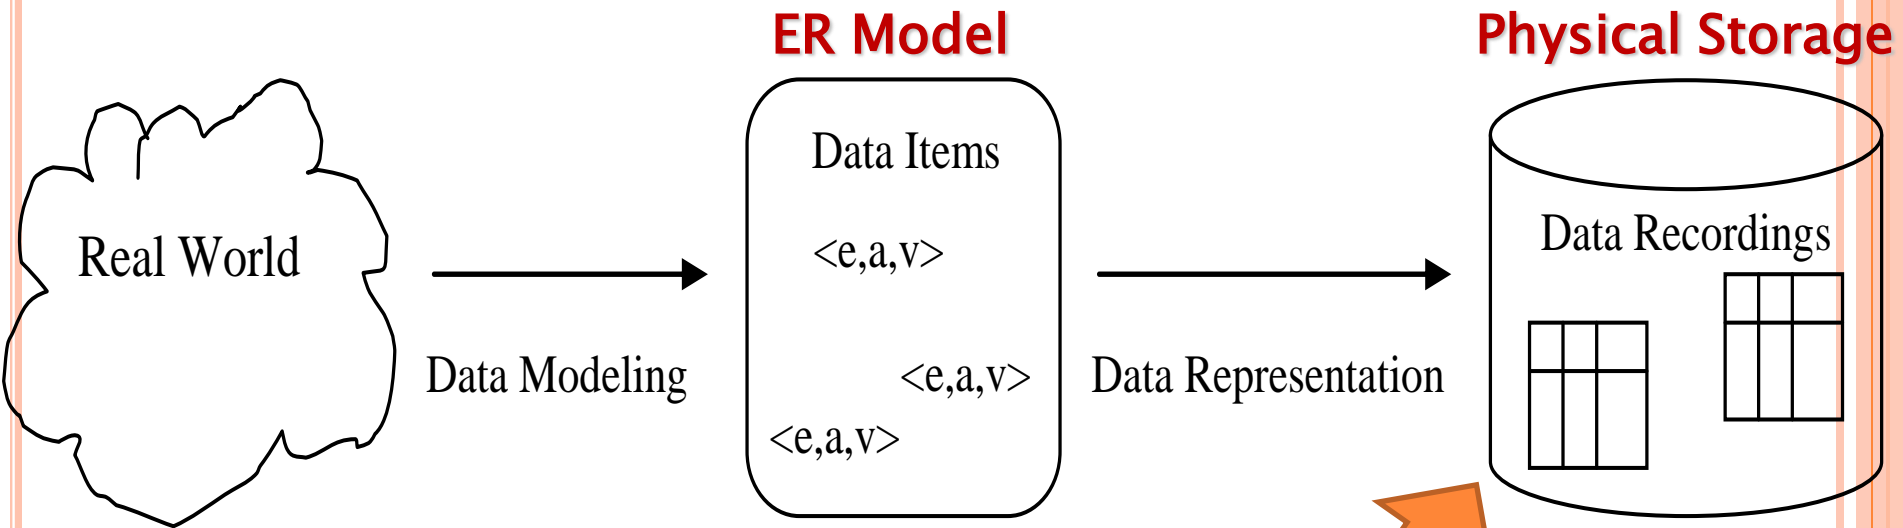
## CHAPTER 8

# OUTLINE

- Overview of Physical Database Design
- Inputs of Physical Database Design
- File Structures
- Query Optimization
- Index Selection
- Additional Choices in Physical Database Design



# DATA MODELING



## Logical Data


StdSSN	StdLastName	StdMajor	StdClass	StdGPA
123-45-6789	WELLS	IS	FR	3.00
124-56-7890	NORBERT	FIN	JR	2.70
234-56-7890	KENDALL	ACCT	JR	3.50

# OVERVIEW OF PHYSICAL DATABASE DESIGN

- Importance of the process and environment of physical database design
  - Process: inputs, outputs, objectives
  - Environment: file structures and query optimization
- Physical Database Design is characterized as a series of *decision-making processes*.
- Decisions involve the storage level of a database: *file structure and optimization choices*.



# STORAGE LEVEL OF DATABASES




- The storage level is closest to the hardware and operating system.
- CPU can process data that are stored in main memory.




# STORAGE LEVEL OF DATABASES

- At the logical level,
  - A database consists of many tables
  - A table consists of many logical records
- At the storage level,
  - A **table** is a **file**
  - A **file** is a collection of **physical records** organized for efficient access.




# RELATIONSHIPS BETWEEN LOGICAL RECORDS (LR) AND PHYSICAL RECORDS (PR)


(a)



(b)



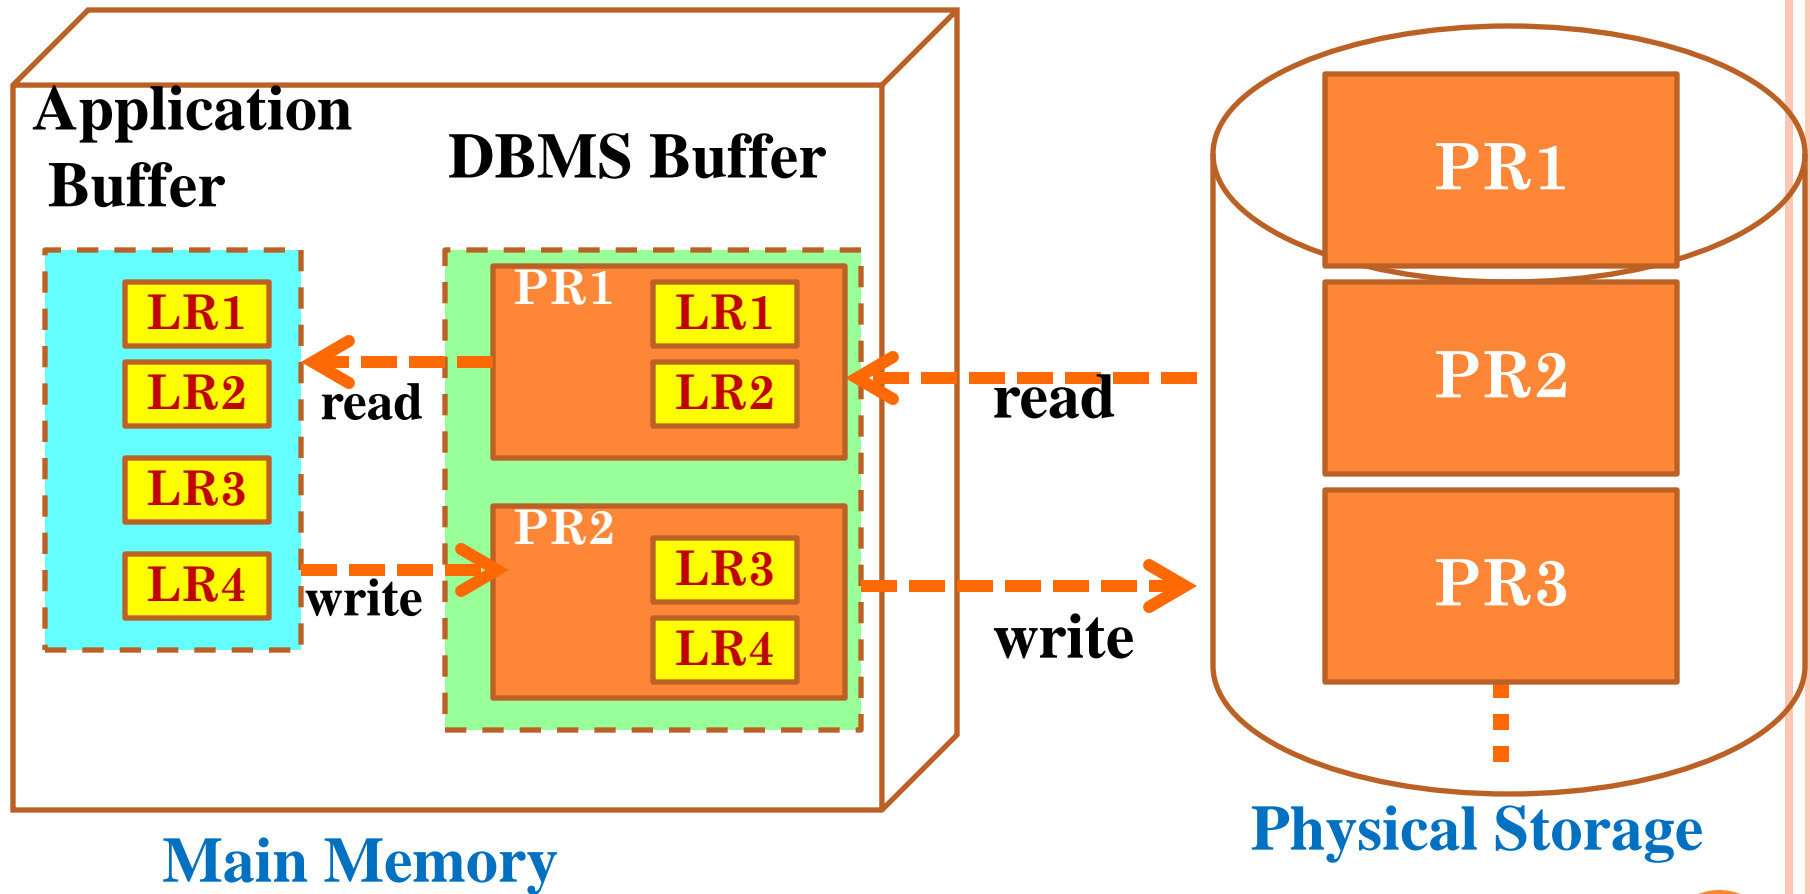
(c)



**A logical record (LR) is a row in a table**



# STORAGE LEVEL OF DATABASES





# OBJECTIVES

- Minimize response time to access and change a database.
- Minimizing computing resources is a substitute measure for response time.
- Database resources
  - Physical record transfers
  - CPU operations
  - Communication network usage (distributed processing)

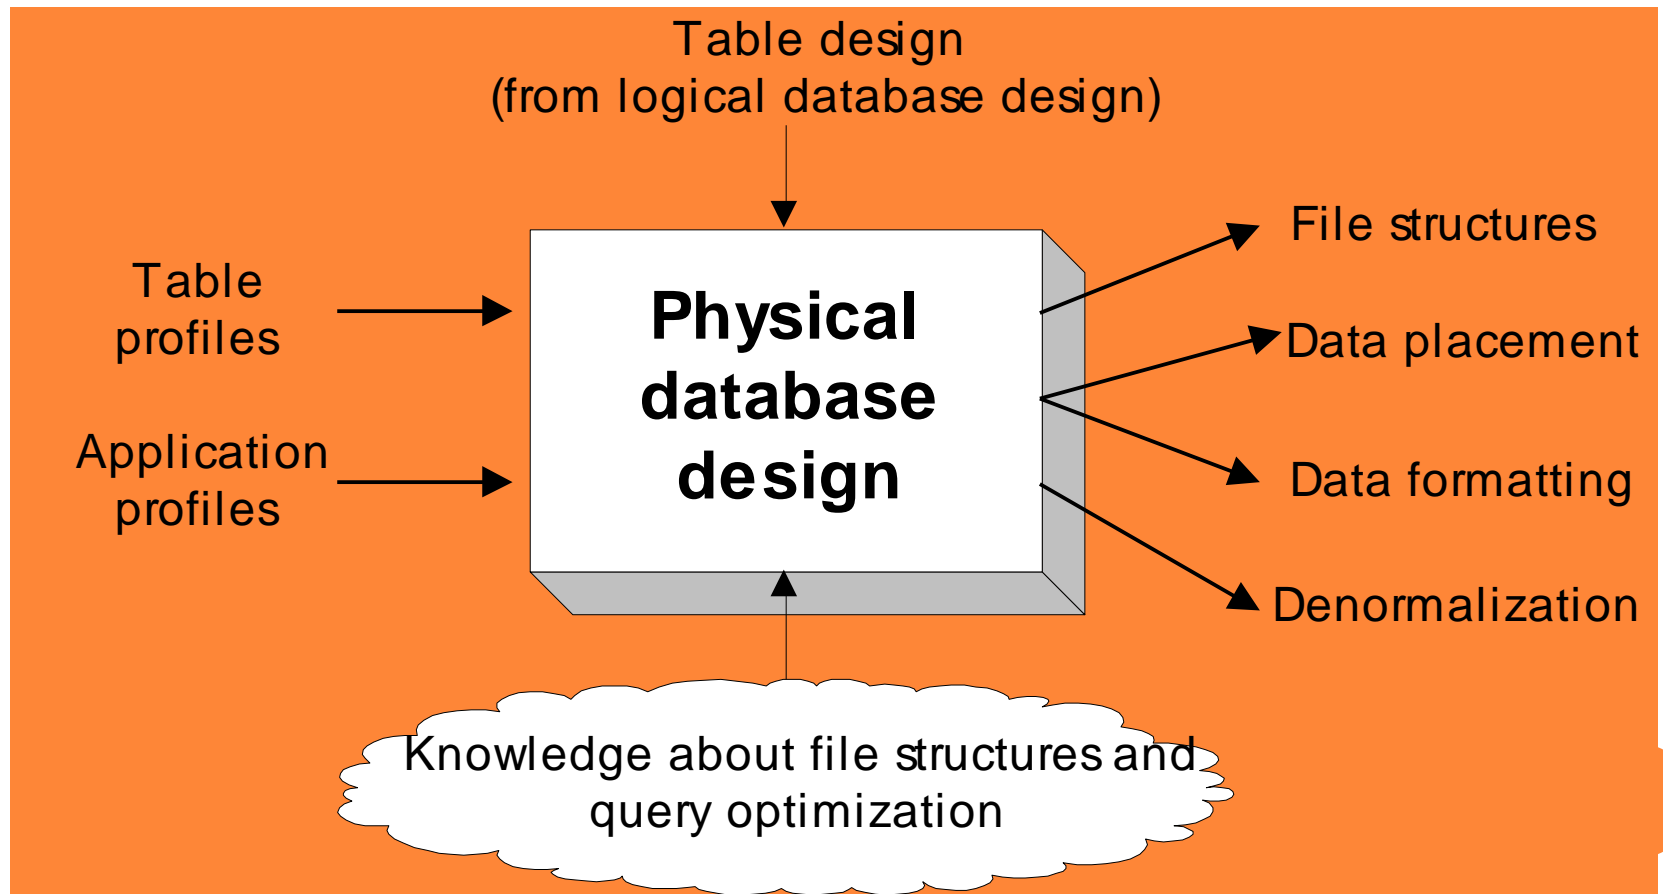


# CONSTRAINTS

- Main memory and disk space are considered as constraints rather than resources to minimize.
- Minimizing main memory and disk space can lead to high response times.
- Thus, reducing the number of physical record accesses can improve response time.
- CPU usage also can be a factor in some database applications.



# INPUTS, OUTPUTS, AND ENVIRONMENT



# INPUTS OF PHYSICAL DATABASE DESIGN

- Physical database design requires inputs specified in sufficient detail.
- Table profiles and application profiles are important and sometimes difficult-to-define inputs.



# TABLE PROFILE

- A table profile summarizes a table as a whole, the columns within a table, and the relationships between tables.

## Typical Components of a Table Profile

Component	Statistics
Table	Number of rows and physical records
Column	Number of unique values, distribution of values
Relationship	Distribution of the number of related rows

# APPLICATION PROFILES

- Application profiles summarize the queries, forms, and reports that access a database.

## Typical Components of an Application Profile

Application Type	Statistics
Query	Frequency; distribution of parameter values
Form	Frequency of insert, update, delete, and retrieval operations to the main form and the subform
Report	Frequency; distribution of parameter values

# FILE STRUCTURES

- Selecting among alternative file structures is one of the most important choices in physical database design.
- In order to choose intelligently, you must understand characteristics of available file structures.



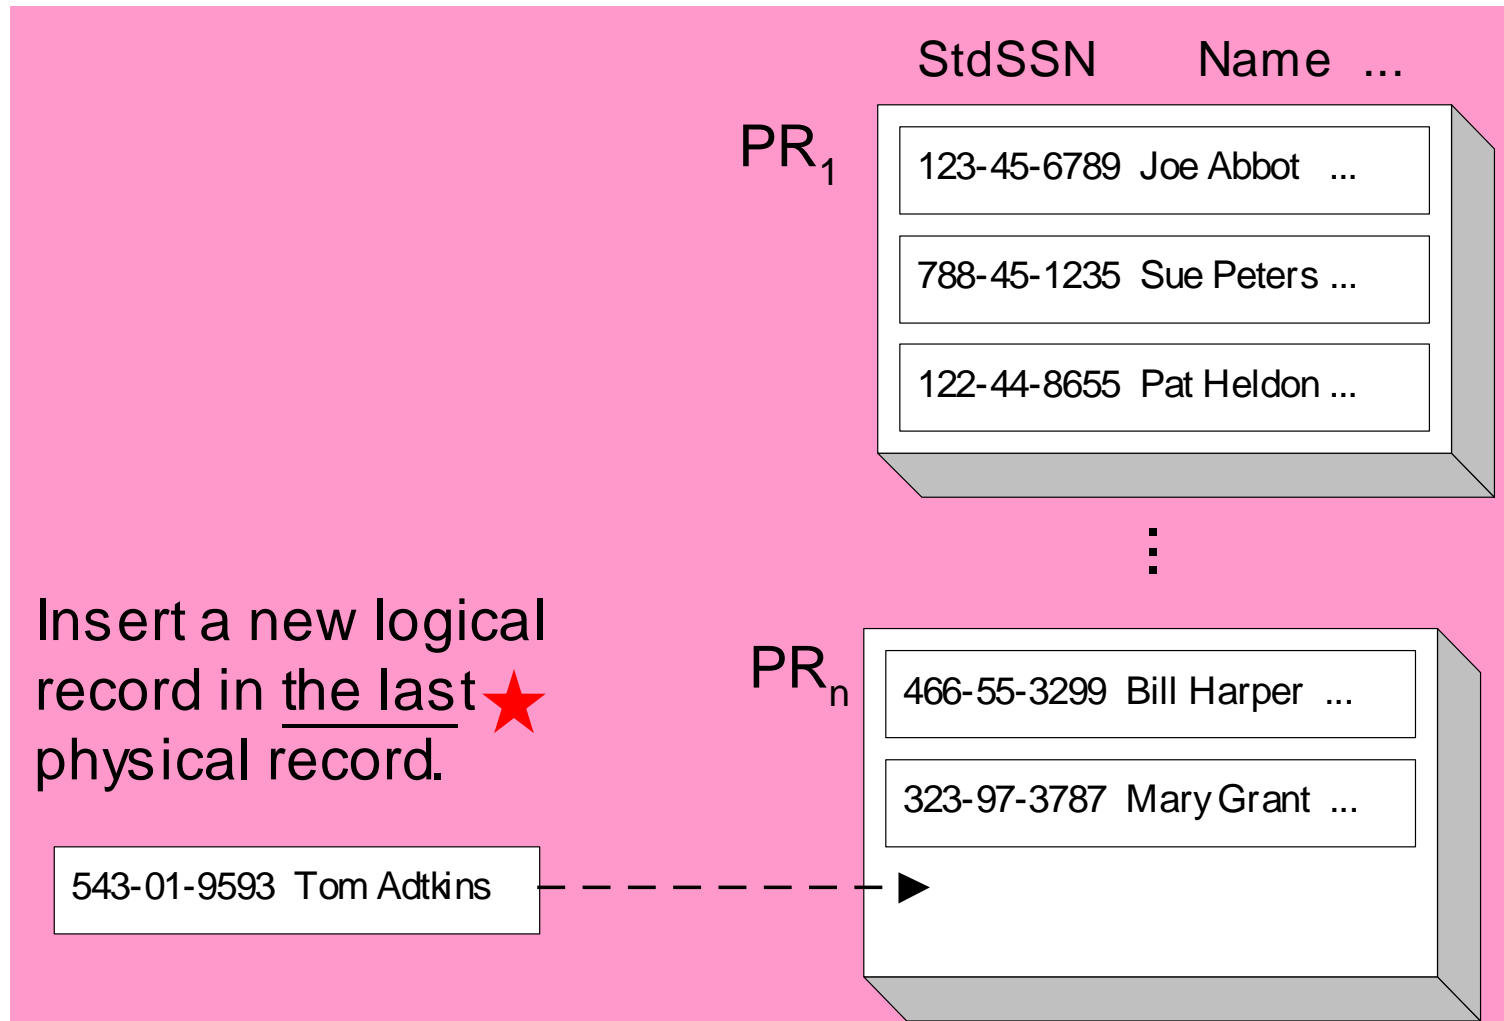
# SEQUENTIAL FILES

- Simplest kind of file structure
- Unordered: insertion order
- Ordered: key order
- Simple to maintain
- Provide good performance for processing large numbers of records





# UNORDERED SEQUENTIAL FILE



# ORDERED SEQUENTIAL FILE



Rearrange physical record  
to insert new logical record.

$PR_1$

StdSSN      Name ...

122-44-8655 Pat Heldon ...

123-45-6789 Joe Abbot ...

323-97-3787 Mary Grant ...


⋮

$PR_n$

466-55-3299 Bill Harper ...

788-45-1235 Sue Peters ...

543-01-9593 Tom Adtkins



# HASH FILES

- Support fast access unique key value
- Converts a key value into a physical record address
- Mod function: typical hash function
  - **Divisor**: large prime number close to the file capacity
  - **Physical record number**: hash function plus the starting physical record number

$$122448655 \bmod 97 = 26,$$


$$26 + \text{starting PR\#} = 176 = \text{location of data}$$

# EXAMPLE: HASH FUNCTION CALCULATIONS FOR STDSSN KEY

**122448655 mod 97 = 26,**

**26+starting PR# = 176**

StdSSN	StdSSN Mod 97	PR Number
<b>122448655</b>	<b>26</b>	<b>176</b>
123456789	39	189
323973787	92	242
466553299	80	230
788451235	24	174
543019593	13	163



# HASH FILE AFTER INSERTIONS

**PR**<sub>163</sub>

543-01-9593 Tom Adtkins

**PR**<sub>189</sub>

123-45-6789 Joe Abbot

⋮

⋮

**PR**<sub>174</sub>

788-45-1235 Sue Peters

**PR**<sub>230</sub>

466-55-3299 Bill Harper

⋮

⋮

**PR**<sub>176</sub>

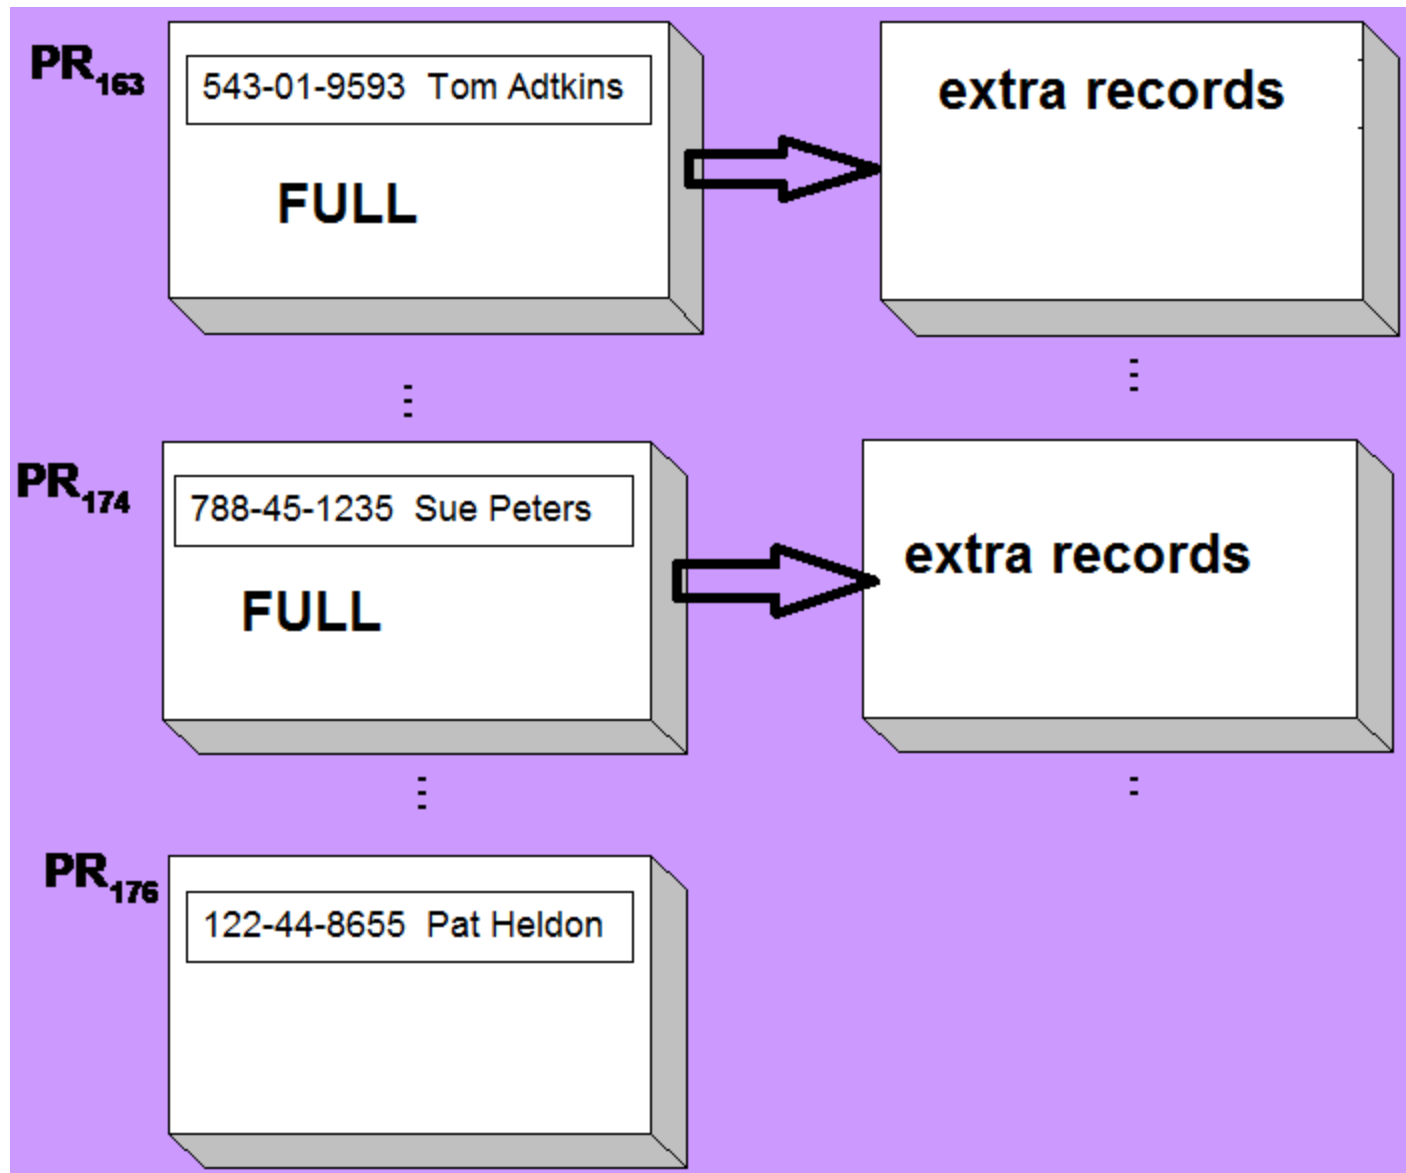
122-44-8655 Pat Heldon

**PR**<sub>242</sub>

323-97-3787 Mary Grant



# HANDLING COLLISIONS



# LINEAR PROBE COLLISION HANDLING DURING AN INSERT OPERATION

Home address = Hash function value + Base address

$(122448946 \bmod 97 = 26) + 150$

...

PR<sub>176</sub>

122-44-8655 Pat Heldon

122-44-8752 Joe Bishop

122-44-8849 Mary Wyatt

Home address (176) is full.

122-44-8946 Tom Atkins

PR<sub>177</sub>

122-44-8753 Bill Hayes

Linear probe to find physical record with space

...



# B TREES





## BTREE CHARACTERISTICS:

### ○ Balanced

- Max access= height of Btree

### ○ Bushy: multi-way tree

- ideal Btree : wide (bushy) but short (few levels).

### ○ Block-oriented

- each node is a physical record

### ○ Dynamic

- Btree changes as logical records are inserted and deleted




# WHY BTREES ?

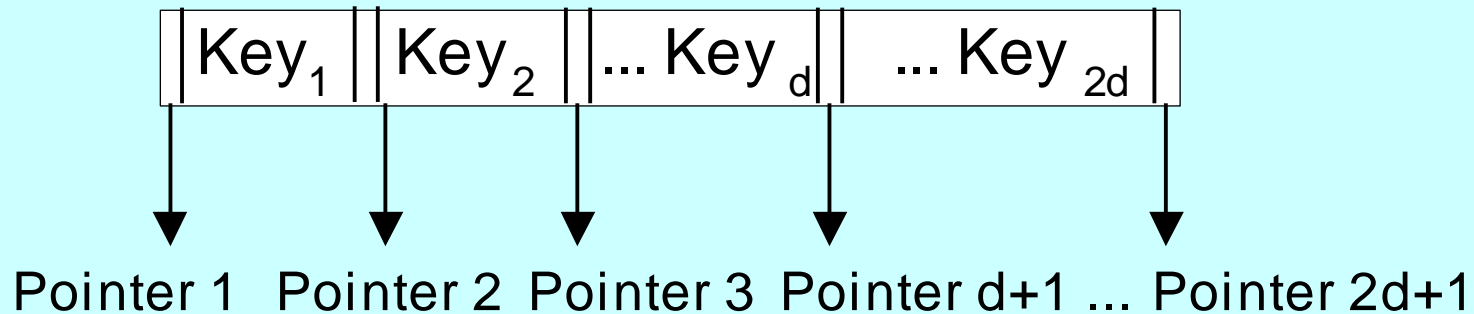
- Sequential files →
  - well on sequential search
  - poorly on key search
- Hash files →
  - well on key search
  - poorly on sequential search,
- Btree is a compromise and widely used file structure.
  - good performance on both sequential search and key search.



# STRUCTURE OF A BTREE OF HEIGHT 3



# BTREE NODE CONTAINING KEYS AND POINTERS




Each non root node contains at least half capacity ( $d$  keys and  $d+1$  pointers).


Each non root node contains at most full capacity ( $2d$  keys and  $2d+1$  pointers).

maximum of two keys


insert 1.




insert 2.



insert 3.




**insert 4.**



**insert 5 ?**




**Split**




insert 5.

combine




**insert 6.**




**7 ?full**

**insert 7.**


**combine**








**After  
insert 7.**




# BTREE INSERTION EXAMPLES


(a) Initial Btree



(b) After inserting 55




(c) After inserting 58




# BTREE DELETION EXAMPLES


(a) Initial Btree




(b) After deleting 60



(c) After deleting 65




**(b) After deleting 60**



**Deleting 65 needs to restructure the tree  
because it will not be half full**



## Deleting 65



**1. Combine nodes**

**22, 28, 35, 45, 50**

**2. Split nodes**



**22, 28**

**35, 45, 50**




**Deleting 65**


**3. Insert 35 in upper node**



# AFTER DELETE 65



**(c) After deleting 65**



# COST OF OPERATIONS

- The **height of Btree** dominates the number of physical record accesses operation.
- Logarithmic search cost
  - Upper bound of height: log function'
  - Log base: minimum number of keys in a node
- The cost to insert a key = [the cost to locate the nearest key] + [the cost to change nodes].  
*or add a new level*






# B+TREE

- Provides improved performance on sequential and range searches.
- In a B+tree, all keys are redundantly stored in the leaf nodes.
- To ensure that physical records are not replaced, the B+tree variation is usually implemented.






**$d_1, d_2, \dots$ : Data Pointers point to physical records**



Index keys:  
support **key** search



Sequence list of all index  
keys


# INDEX MATCHING

- Determining usage of an index for a query
- Complexity of condition determines match.
- Single column indexes: =, <, >, <=, >=, IN <list of values>, BETWEEN, IS NULL, LIKE 'Pattern' (meta character not the first symbol)
- Composite indexes: more complex and restrictive rules



# BTREE AND HASH FILES

- Work best for columns with unique values
  - E.g., ID
- Btrees index nodes can store a list of row identifiers for non unique columns
  - the list of row identifiers can be very long



# BITMAP INDEX

- Can be useful for stable columns with few values
- Bitmap:
  - String of bits: **0 (no match) or 1 (match)**
  - One bit for each row
- Bitmap index record
  - Column value (**non unique columns**)
  - Bitmap (e.g., 0011100)
  - DBMS converts bit position into row identifier.



# BITMAP INDEX ON STUDENT GRADE

Student Id	Student Grade				
	A	B	C	D	F
101	0	1	0	0	0
102	0	1	0	0	0
103	1	0	0	0	0
104	0	0	1	0	0
105	0	1	0	0	0
106	0	0	0	1	0
107	0	1	0	0	0
108	0	1	0	0	0
109	1	0	0	0	0
110	0	1	0	0	0
111	0	0	1	0	0
112	0	1	0	0	0
113	1	0	0	0	0
114	0	1	0	0	0
115	0	0	0	1	0
116	1	0	0	0	0
117	0	0	0	0	1
118	0	1	0	0	0
119	0	1	0	0	0
120	0	0	1	0	0

**0 (no match)**

**1 (match)**



## Faculty Table

## BITMAP INDEX EXAMPLE

RowId	FacSSN	...	FacRank
1	098-55-1234		Asst
2	123-45-6789		Asst
3	456-89-1243		Assc
4	111-09-0245		Prof
5	931-99-2034		Asst
6	998-00-1245		Prof
7	287-44-3341		Assc
8	230-21-9432		Asst
9	321-44-5588		Prof
10	443-22-3356		Assc
11	559-87-3211		Prof

### Bitmap Index on FacRank

12	FacRank	Bitmap
	<b>Asst</b>	<b>110010010001</b>
	<b>Assoc</b>	<b>001000100100</b>
	<b>Prof</b>	<b>000101001010</b>



# Faculty Table

RowId	FacSSN	...	FacRank	
1	098-55-1234		Asst	<b>1</b> ← <b>Asst</b>
2	123-45-6789		Asst	<b>1</b> ← <b>Asst</b>
3	456-89-1243		Assc	<b>0</b>
4	111-09-0245		Prof	<b>0</b>
5	931-99-2034		Asst	<b>1</b> ← <b>Asst</b>
6	998-00-1245		Prof	<b>0</b>
7	287-44-3341		Assc	<b>0</b>
8	230-21-9432		Asst	<b>1</b> ← <b>Asst</b>
9	321-44-5588		Prof	<b>0</b>
10	443-22-3356		Assc	<b>0</b>
11	559-87-3211		Prof	<b>0</b>
12	220-44-5688		Asst	<b>1</b> ← <b>Asst</b>



FacRank	Bitmap
<b>Asst</b>	<b>110010010001</b>
<b>Assoc</b>	<b>001000100100</b>
<b>Prof</b>	<b>000101001010</b>

**Bitmap Index on FacRank**

Column value

Bitmap

FacRank	Bitmap
Asst	110010010001
Assoc	001000100100
Prof	000101001010

- Asst , 110010010001
- Row# 1,2,5,8,12 are Asst. Prof.



# BITMAP JOIN INDEX

- Bitmap identifies rows of a related table.
- Represents a pre-computed join
- Typically used in query dominated environments such as data warehouses (Chapter 16)

CourseNo	BitmapJoin
204351	110
204111	001



CourseNo	Bitmap
204351	110
204111	?

Course

CourseNo	CrsDesc	CrsUnits
204351	Database	3
204111	C#	3

Primary key

Offering

OfferNo	OffLocation	OffTime	CourseNo
111	CPE 203	1/2014	204351
222	CPE 204	2/2014	204351
333	CPE 204	2/2014	204111

Foriegn key



# SUMMARY OF FILE STRUCTURES

	<b>sequential files</b>				
	<b>Unordered</b>	<b>Ordered</b>	<b>Hash</b>	<b>B+tree</b>	<b>Bitmap</b>
<b>Sequential search</b>	Y	Y	Extra PRs	Y	N
<b>Key search</b>	Linear	Linear	Constant time	Logarithmic	Y
<b>Range search</b>	N	Y	N	Y	Y
<b>Usage</b>	Primary only	Primary only	Primary or secondary	Primary or secondary	Secondary only

### 3.A BITMAP INDEX: RANGE SEARCHES

CourseNo		Bitmap	
204101	} 101-102	1100001	} Union = 1110011
204102		0010010	
204103		0001000	



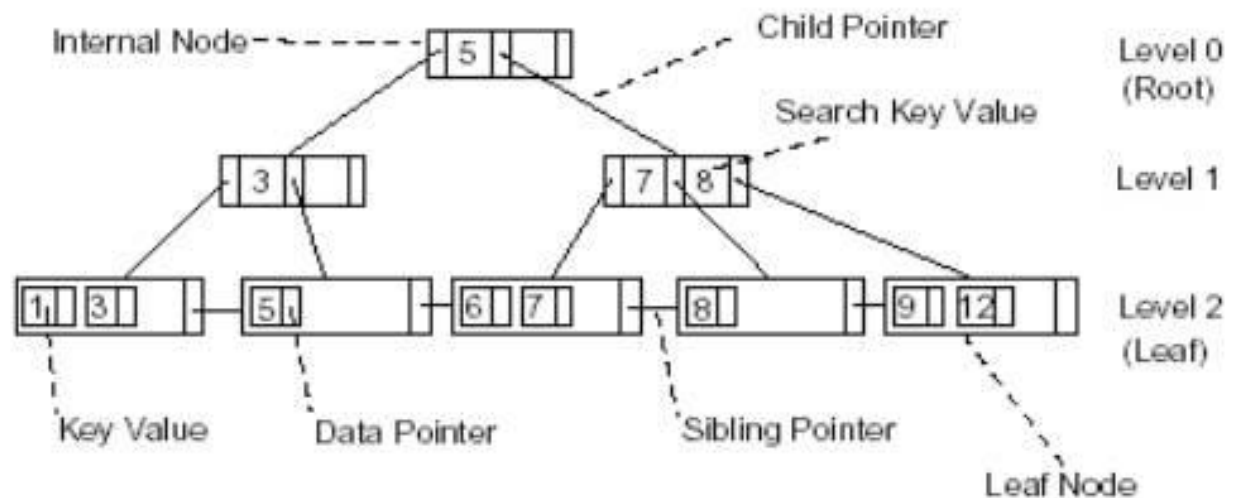
## 4.PRIMARY FILE STRUCTURE VS SECONDARY FILE STRUCTURE

### ○ Primary file structure

- store all the data of a table

### ○ Secondary file structure


- store only key data along with pointers to the data records



## sequential files



	Unordered	Ordered	Hash	B+tree	Bitmap
Usage	Primary only	Primary only	Primary or secondary	Primary or secondary	Secondary only



PR<sub>176</sub>

122-44-8655 Pat Heldon

122-44-8752 Joe Bishop

122-44-8849 Mary Wyatt

**Primary Data**

PR<sub>177</sub>

122-44-8753 Bill Hayes

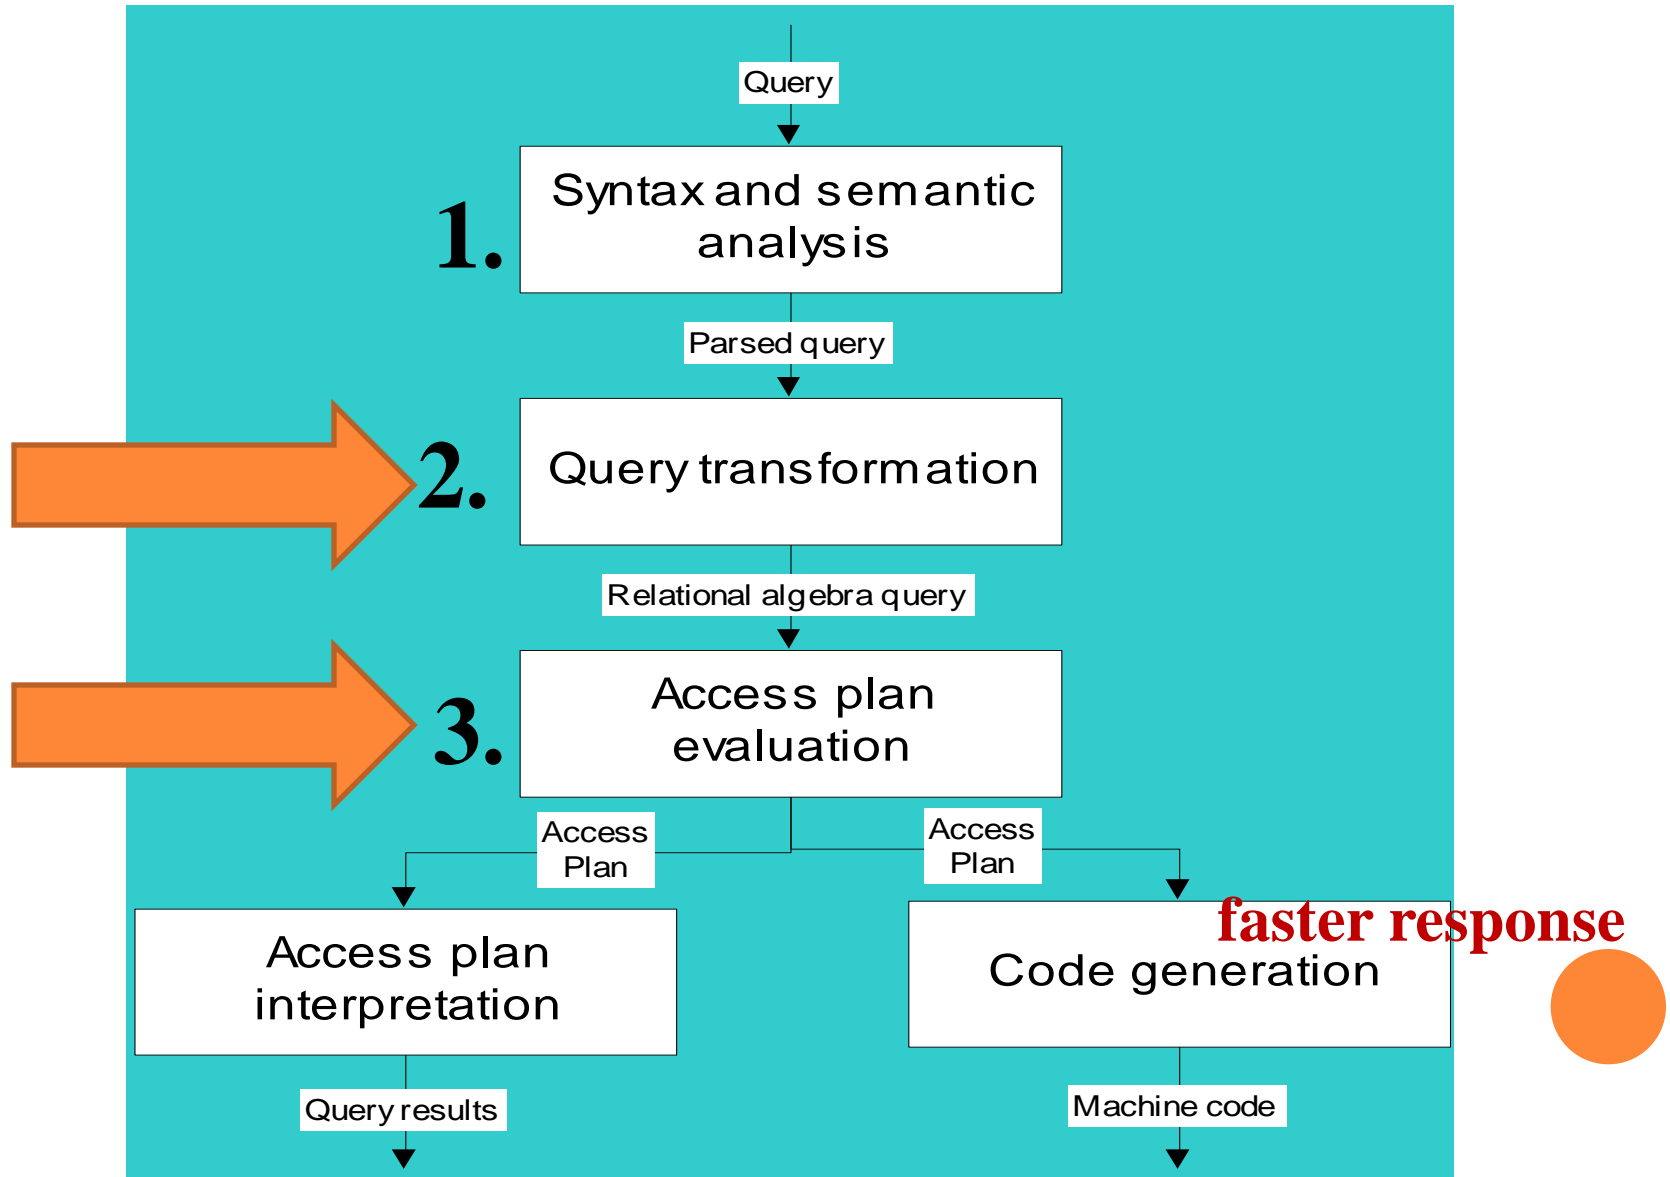


# QUERY OPTIMIZATION

- Query optimizer determines implementation of queries.
- Major improvement in software productivity
- You can sometimes improve the optimization result through knowledge of the optimization process.



# TRANSLATION TASKS



# TRANSLATION TASKS

1. Analyzes a query for syntax and simple semantic errors
2. Transforms a query into a simplified and **standardized format** so that the query can be executed faster.
3. Determines how to implement an access plan
  - **Access plan:** how to implement a query as operations on files
    - file structures to access tables,
    - the order of joining tables,
    - the algorithm to join tables

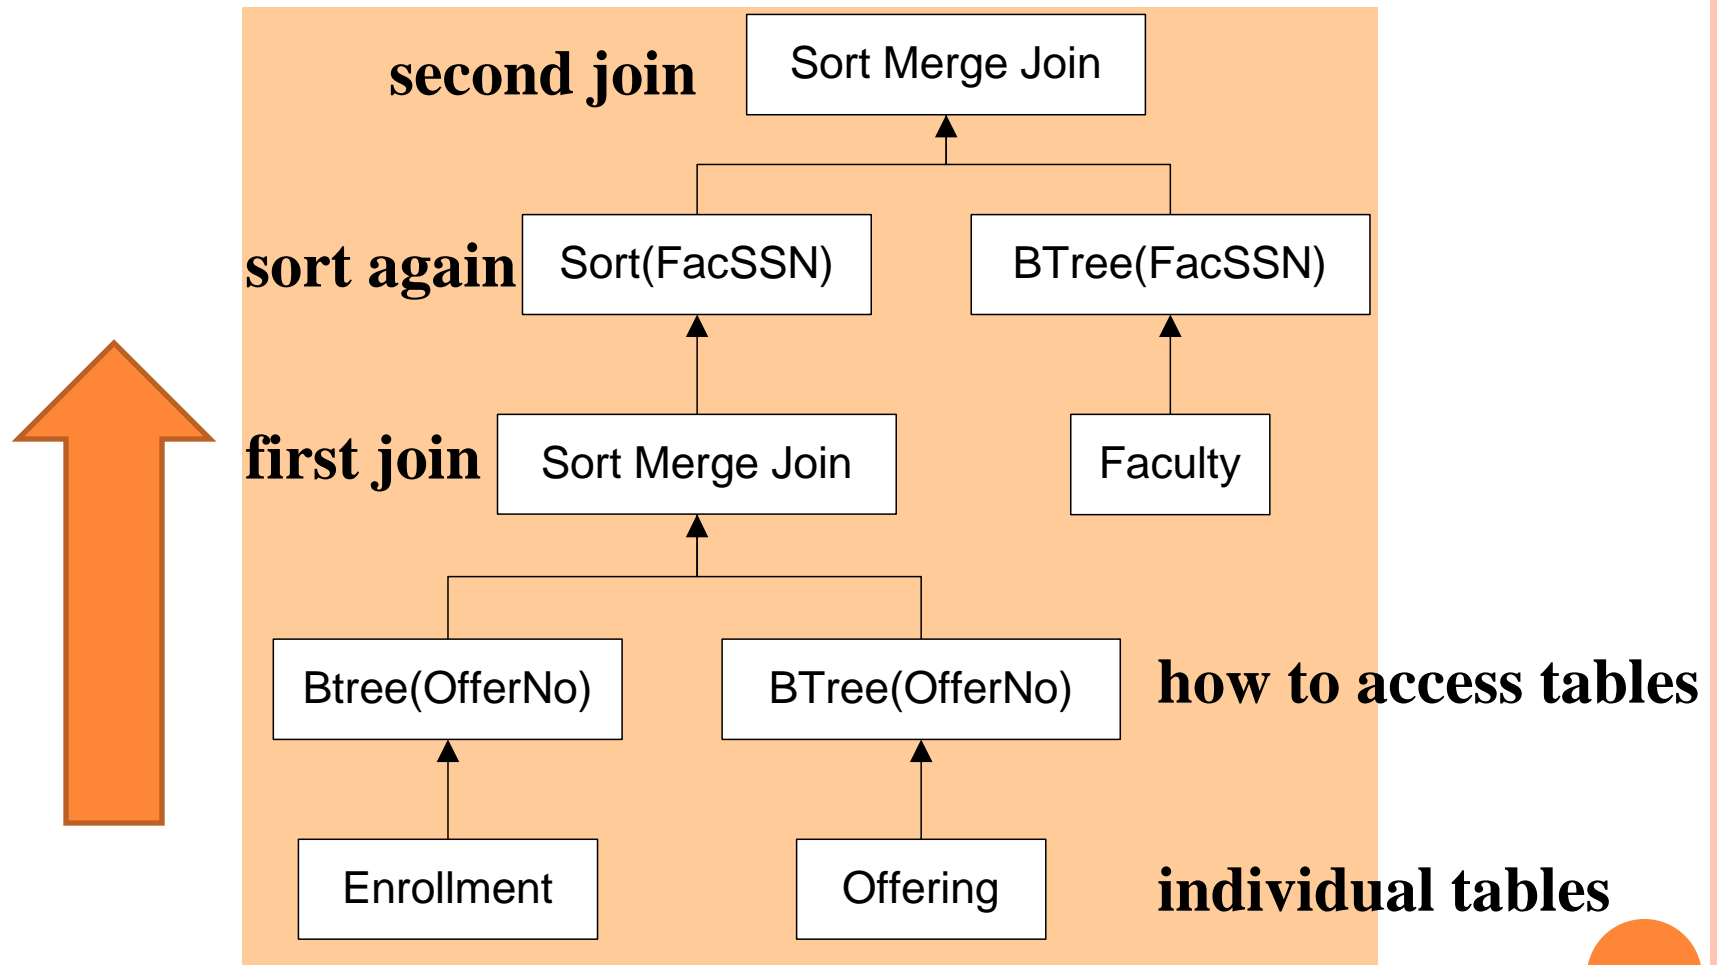


# ACCESS PLAN

- Each operation in an access plan has a **corresponding cost formula** that estimates the physical record accesses and CPU operations.
- The cost formulas use table profiles to estimate the **number of rows** in a result.
- The query optimization component **chooses the access plan with the lowest cost.**



# ACCESS PLANS



see Oracle Plan in

[http://docs.oracle.com/cd/B19306\\_01/server.102/b14211/ex\\_plan.htm](http://docs.oracle.com/cd/B19306_01/server.102/b14211/ex_plan.htm)

# ACCESS PLAN EVALUATION

- Optimizer evaluates thousands of access plans
- Access plans vary by join order, file structures, and join algorithm.
- Some optimizers can use multiple indexes on the same table.
- Access plan evaluation can consume significant resources
  - ex., when the query contains more than four tables



# OPTIMIZATION TIPS I

- Detailed and current statistics needed
- Save access plans for repetitive queries
- Review access plans to determine problems
- Use hints carefully to improve results



## OPTIMIZATION TIPS II

- Replace Type II nested queries with separate queries.
- For conditions on join columns, test the condition on the parent table.
- Do not use the HAVING clause for row conditions.



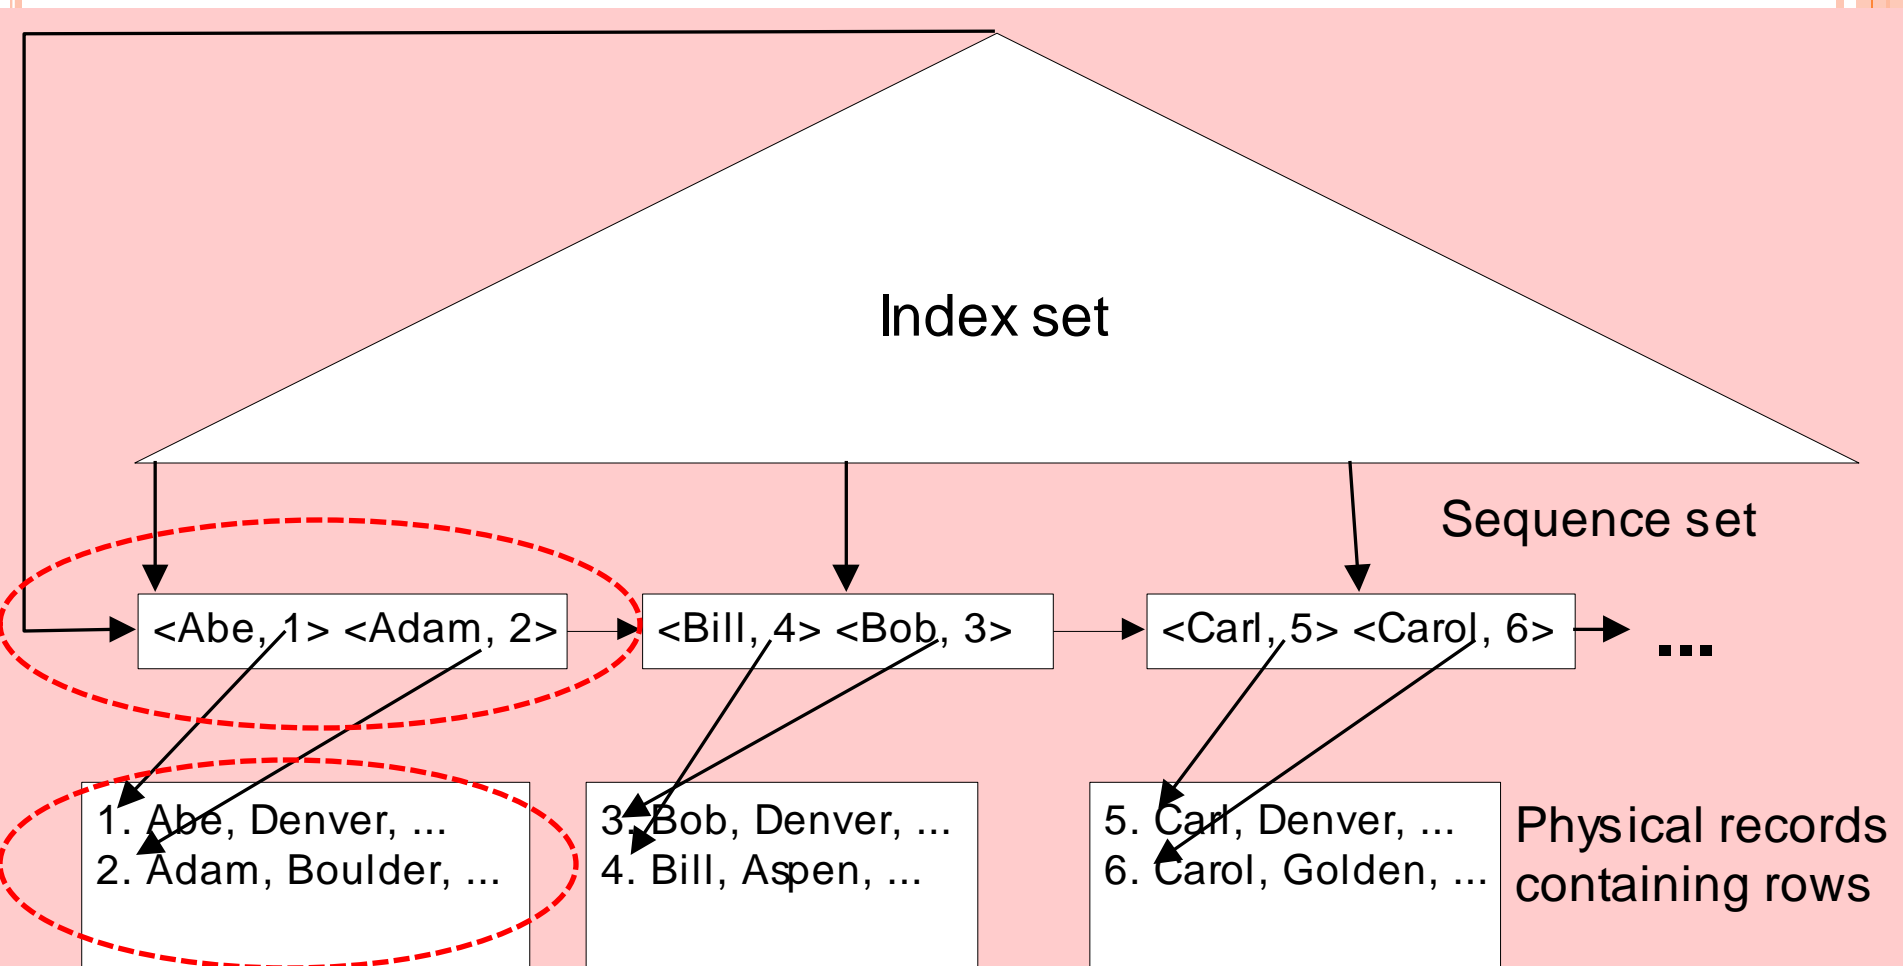


# INDEX SELECTION

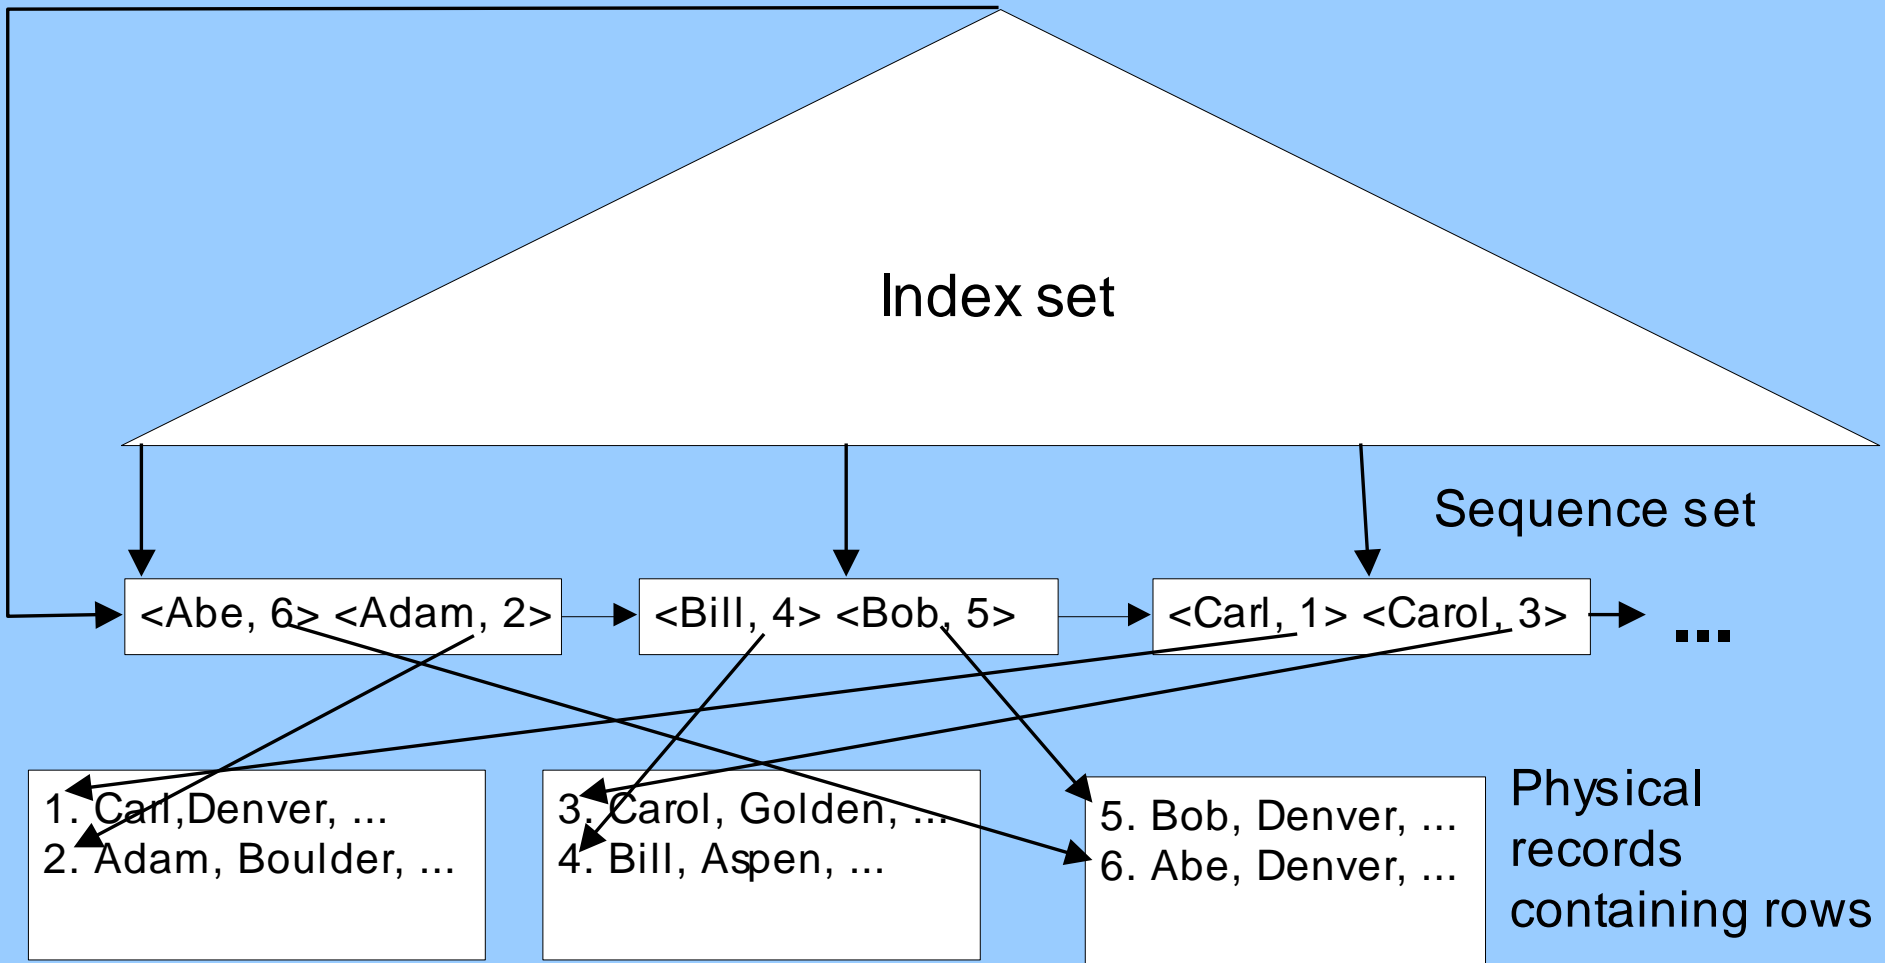
- Most important decision
- Difficult decision
- Choice of clustered and nonclustered indexes




# CLUSTERING INDEX EXAMPLE



# NONCLUSTERING INDEX EXAMPLE



# INPUTS AND OUTPUTS OF INDEX SELECTION



**W: frequency of a statement + its importance**



# TRADE-OFFS IN INDEX SELECTION

- Balance retrieval against update performance
- Nonclustering index usage:
  - Few rows satisfy the condition in the query
  - Join column usage if a small number of rows result in child table
- Clustering index usage:
  - Larger number of rows satisfy a condition than for nonclustering index
  - Use in sort merge join algorithm to avoid sorting
  - More expensive to maintain



# INDEX SELECTION RULES

- A primary key is a good candidate for a clustering index.
- To support joins, consider indexes on foreign keys.
- A frequently updated column is not a good index candidate.
- Volatile tables (lots of insertions and deletions) should not have many indexes.
- Stable columns with few values are good candidates for bitmap indexes if the columns appear in WHERE conditions.



# INDEX CREATION

- To create the indexes, the CREATE INDEX statement can be used.
- The word following the INDEX keyword is the name of the index.
- CREATE INDEX is not part of SQL:1999.

## Example:

```
CREATE INDEX StdGPAIndex ON Student (StdGPA)
CREATE UNIQUE INDEX OfferNoIndex ON Offering
(OfferNo)
CREATE BITMAP INDEX OffYearIndex ON Offering
(OffYear)
```

# DENORMALIZATION

- Additional choice in physical database design
- Denormalization combines tables so that they are easier to query.
- Use carefully because normalized designs have important advantages.





# NORMALIZED DESIGNS

- Better update performance
- Require less coding to enforce integrity constraints
- Support more indexes to improve query performance

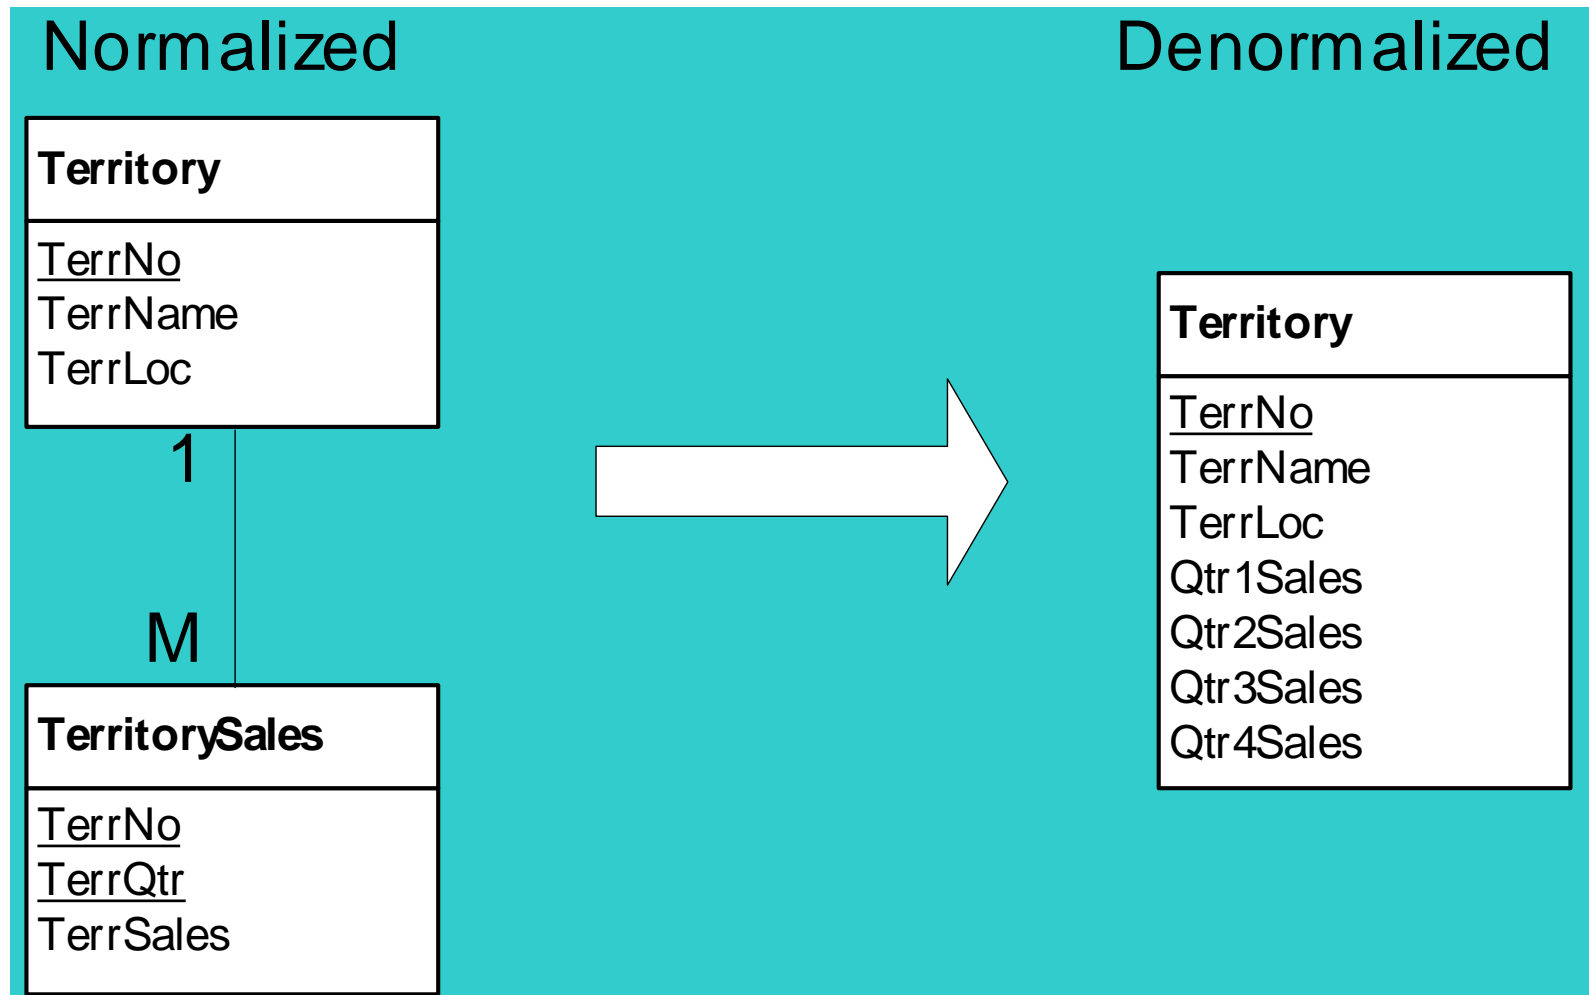


# REPEATING GROUPS

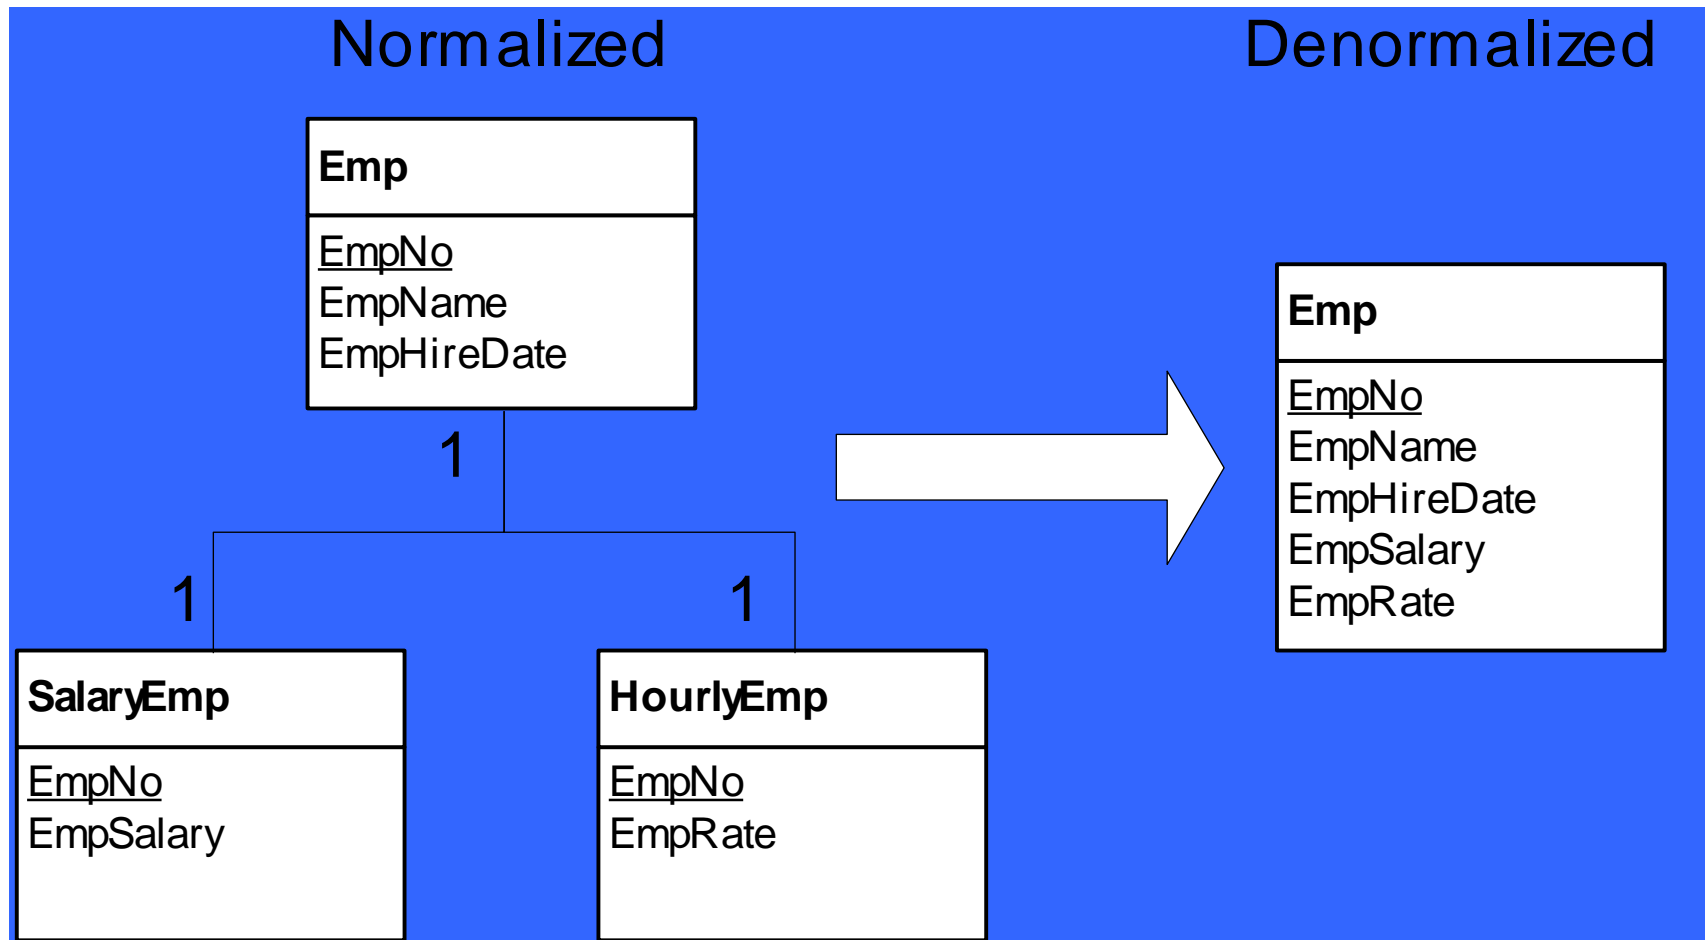
- A repeating group is a collection of associated values.
- The rules of normalization force repeating groups to be stored in an M table separate from an associated one table.
- If a repeating group is always accessed with its associated one table, denormalization may be a reasonable alternative.



# DENORMALIZING A REPEATING GROUP





# DENORMALIZING A GENERALIZATION HIERARCHY




# CODES AND MEANINGS

## Normalized



## Denormalized

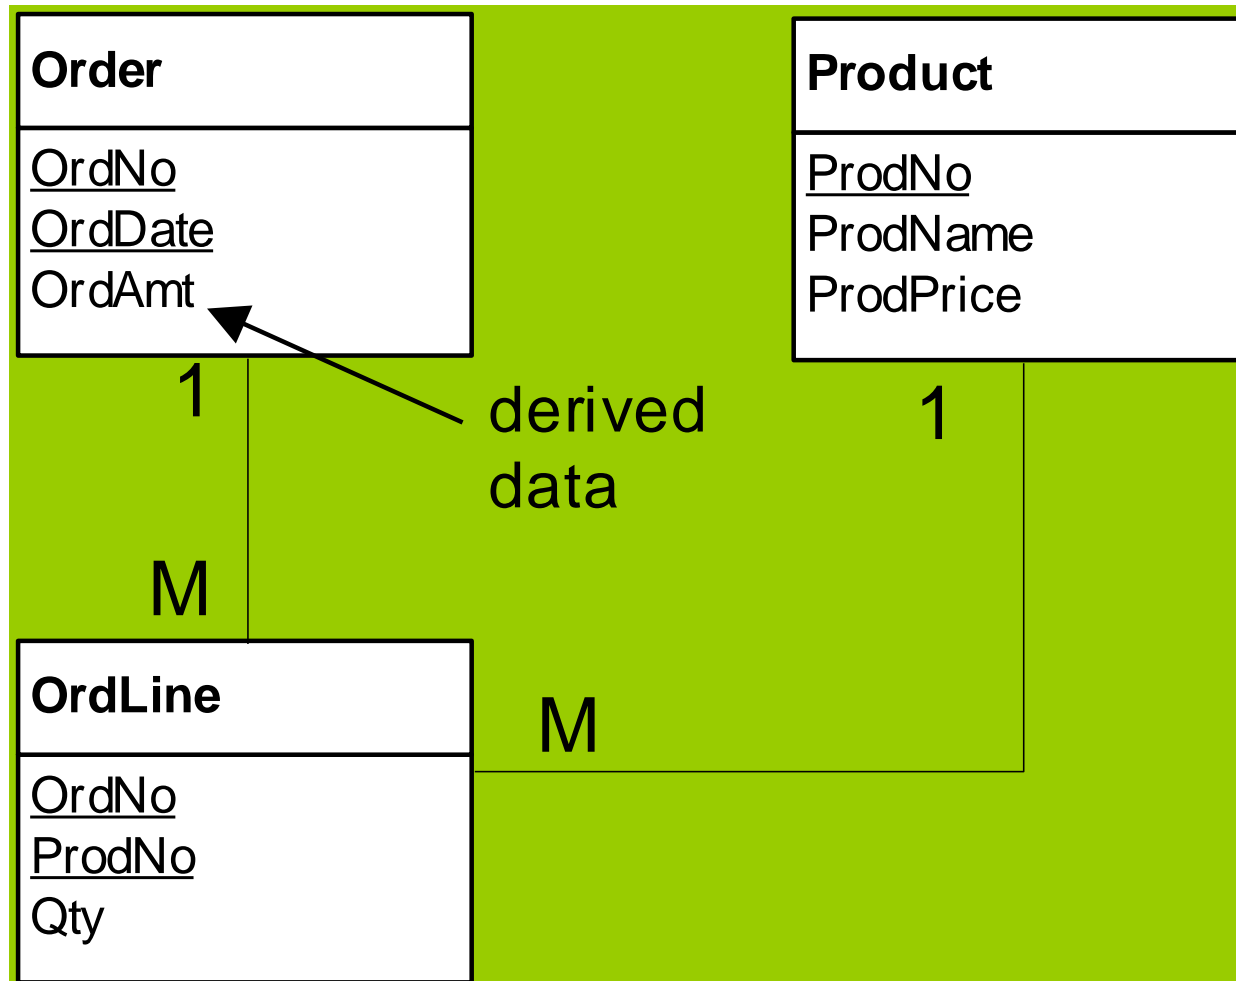


# RECORD FORMATTING

- **Compression** is a trade-off between input-output and processing effort.
- **Derived data** is a trade-offs between query and update operations.



# STORING DERIVED DATA TO IMPROVE QUERY PERFORMANCE



# PARALLEL PROCESSING

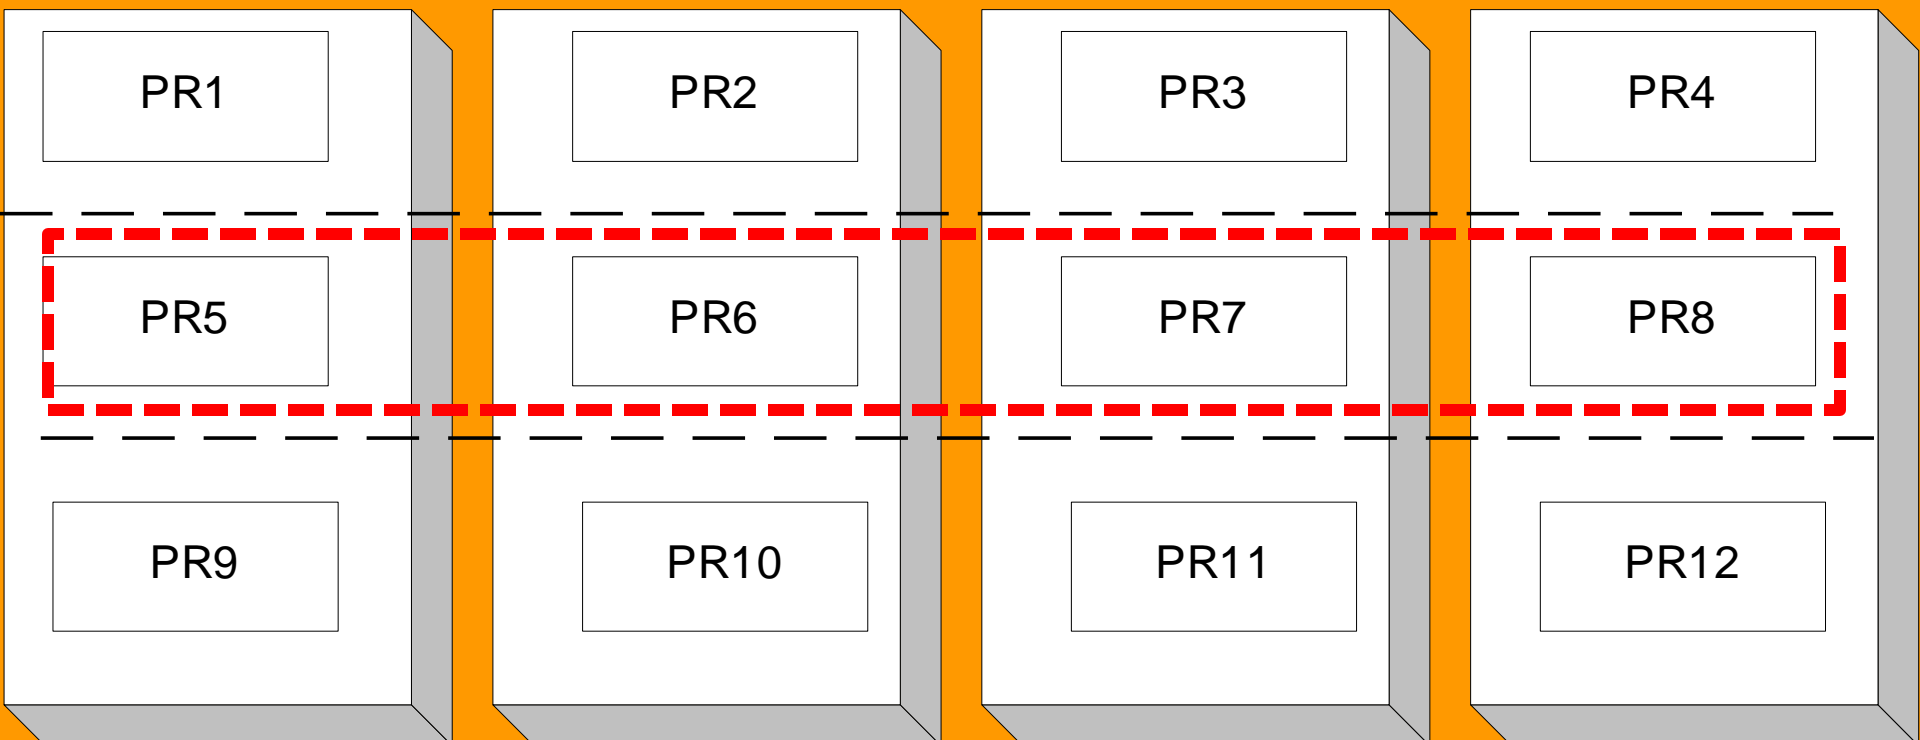
- Retrieving many records can be improved by **reading physical records in parallel.**
- Many DBMSs provide parallel processing capabilities with RAID systems.
- RAID is a collection of disks (a disk array) that operates as a single disk.





# STRIPING IN RAID STORAGE SYSTEMS

Each stripe consists of four adjacent physical records. Three stripes are shown separated by dotted lines.



## OTHER WAYS TO IMPROVE PERFORMANCE

- Transaction processing: add computing capacity and improve transaction design.
- Data warehouses: add computing capacity and store derived data.
- Distributed databases: allocate processing and data to various computing locations.



# SUMMARY

- Goal: minimize computing resources
- Table profiles and application profiles must be specified in sufficient detail.
- Environment: file structures and query optimization
- Monitor and possibly improve query optimization results
- Index selection: most important decision
- Other techniques: denormalization, record formatting, and parallel processing

