

Transaction Management



Transaction ?



Recovery process?

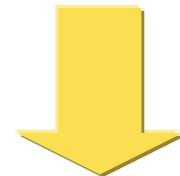
Failure?



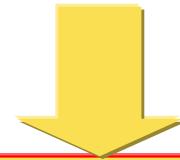
Withdraw Money



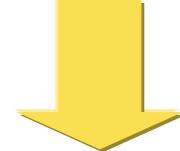
Enter PWD



Enter account Info



Withdraw 5000



Failure



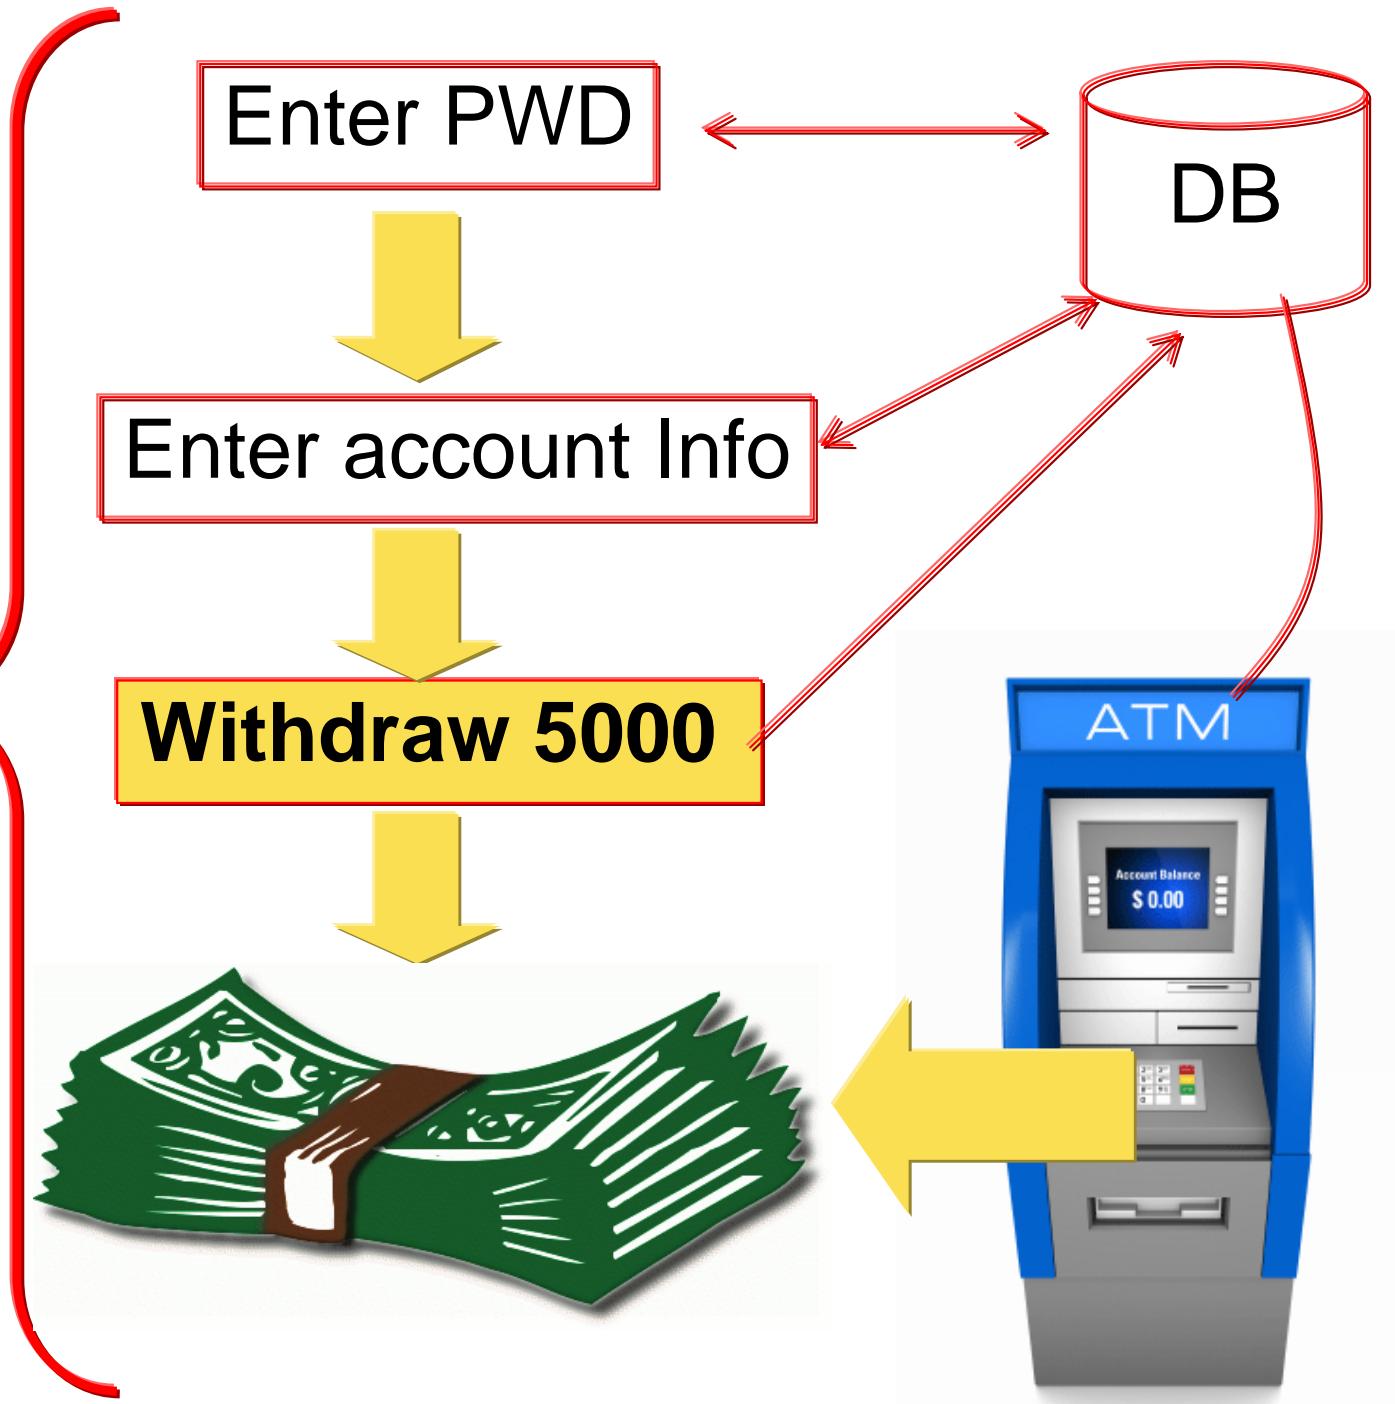
Balance updated ?



© Can Stock Photo - csp3931903



Transaction



Outline

- ✿ Transaction basics
- ✿ Concurrency control
- ✿ Recovery management
- ✿ Transaction design issues
- ✿ Workflow management



Transaction Definition

- ✿ Supports daily operations of an organization
- ✿ Become more important with the growth of the internet



A STAR ALLIANCE MEMBER



Booking.com



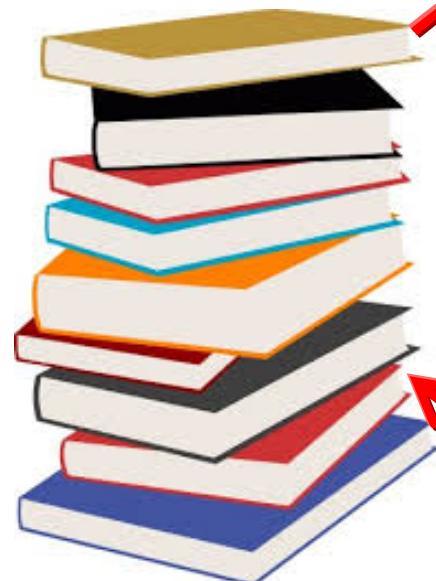
Transaction Definition

- ✿ Collection of database operations
- ✿ Reliably and efficiently processed as one unit of work
- ✿ No lost data
 - ✿ Interference among multiple users
 - ✿ Failures (operating system, program, disk, ...)



Library Example

10 copies of
a database book



1.borrow



3.return

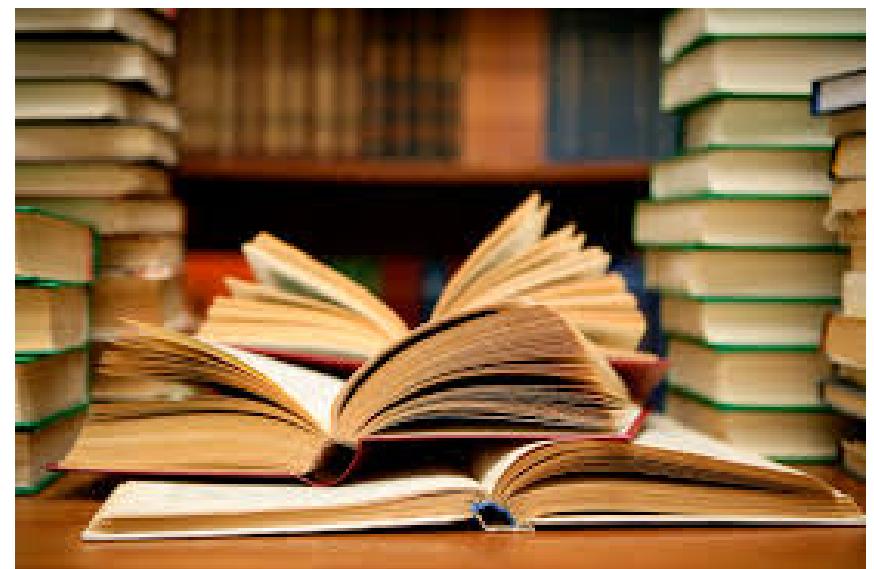
2.Update

4.Update all



Library Example

- ✿ Borrow for read only
 - ✿ concurrently
- ✿ Borrow for update
 - ✿ one person at a time



Airline Transaction Example



Home Flights Vacation packages

NEW FLIGHT SEARCH

Leaving from: Nearby cities
Toronto, ON (YTO-All Airports)

Going to: Nearby cities
New York, NY, United States (NYC-A)

Departing: 16/7/2012 Returning: 20/7/2012

[More options](#)

1. Departure & return cities, dates

To: YO (YTO) to New York, NY, United States

2, 2 travellers

[Show flight summary grid](#)

1. Choose Your Departing Flight

Note: Prices are per person for round trip travel, they are e-ticket prices and include all flight taxes and fees. Prices may not include baggage fees or other fees charged directly by the airline.

These results cover a large area with several airports. Please review your choices carefully.

2. Show available flights

Sort by:

REFINE RESULTS

Displaying all results

[Flight Times](#)

Outbound to New York (NYC)

Depart Arrive

Depart 6:15am – 9:40pm

Toronto **YYZ 3:15pm → LGA 4:38pm** New York
AIR CANADA Air Canada 716 Nonstop 1h 23m Return from **C\$289**
avg/person (from C\$578 total)
includes taxes & fees

[Show Flight Details](#)

New SQL statements:

- **START TRANSACTION**
 - defines **beginning of transaction statements**
 - some **DBMSs** omit
- **COMMIT: end of transaction**
- **ROLLBACK**
 - ✿ Undo command
 - ✿ If any error occurs, all changes are deleted from the database
- ✿ **On Error: part of exception handling**



Airline Transaction Example

START TRANSACTION

Display greeting

Get reservation preferences from user

SELECT departure and return flight records

If reservation is acceptable then

UPDATE seats remaining of departure flight record

UPDATE seats remaining of return flight record

INSERT reservation record

Print ticket if requested

End If

On Error: ROLLBACK

COMMIT



If fail?



ATM Transaction Example

START TRANSACTION

Display greeting

If fail?

Get account number,pin,type,amount

SELECT account number,type,balance

If **balance is sufficient** then

UPDATE account by posting debit

INSERT history record

Display message and **dispense cash**

UPDATE history record

Print receipt if requested

End If

On Error: **ROLLBACK**

COMMIT



Transaction Properties (ACID)

- ✿ **Atomic**
- ✿ **Consistent**
- ✿ **Isolated**
- ✿ **Durable**



Transaction Properties (ACID)

✿ **Atomic:** all or nothing

✿ all changes are made or none are made

✿ **Consistent:** database must be consistent before and after a transaction

✿ by the time any transaction ends, each and every reference in the database must be valid.

example



Transaction Properties (ACID)

✿ Enforce referential integrity:

✿ If a transaction consisted of an attempt to delete a record referenced by another, each of the followings would maintain consistency:

- **abort** the transaction, rolling back to the consistent, prior state; or,
- **delete ALL records** that reference the deleted record (*cascade delete*); or,
- **nullify the relevant fields** in **ALL records** that point to the deleted record.

propagation
constraints



OfferNo	CourseNo	OffTerm	OffYear	OffLocation	OffTime	FacSSN	OffDays
1111	IS320	SUMMER	2006	BLM302	10:30:00		MW
1234	IS320	FALL	2005	BLM302			MW
2222	IS460	SUMMER	2005	BLM412			TH
3333	IS320	SPRING	2006	BLM214			MW
4321	IS320	FALL	2005	BLM214			TH
4444	IS320	WINTER	2006	BLM302	10:30:00	543-21-0987	TTH
5555	FIN300	WINTER	2006	BLM207	09:00:00	765-43-2109	MW
5678	IS480	WINTER	2006	BLM202	09:00:00	007-15-1221	MW
5679	IS480	SPRING	2006	BLM202	09:00:00	007-15-1221	MW
6666	FIN450	WINTER	2006	BLM202	09:00:00	007-15-1221	MW
7777	FIN480	SPRING	2006	BLM202	09:00:00	007-15-1221	MW
8888	IS320	SUMMER	2006	BLM202	09:00:00	007-15-1221	MW
9876	IS460	SPRING	2006	BLM202	09:00:00	007-15-1221	MW

Offering

Course

Transaction:
delete IS320

CourseNo	CrsDesc	CrsUnits
FIN300	FUNDAMENTALS OF FINANCE	4
FIN450	PRINCIPLES OF INVESTMENTS	4
FIN480	CORPORATE FINANCE	4
IS320	FUNDAMENTALS OF BUSINESS PROGRAMMING	4
IS460	SYSTEMS ANALYSIS	4
IS470	BUSINESS DATA COMMUNICATIONS	4
IS480	FUNDAMENTALS OF DATABASE MANAGEMENT	4

Transaction Properties(ACID)

- ✿ **Isolated:** no unwanted interference from other *users*
 - ✿ requirement that other operations cannot access data that has been modified during a transaction that has not yet completed
 - ✿ In case of concurrent transactions



Transaction Properties(ACID)

✿ Durable:

✿ the ability of the DBMS to **recover** the **committed transaction** updates against any kind of system failure (HW or SW)

- the data changes will survive system failure, and
- all integrity constraints have been satisfied so the DBMS won't need to reverse the transaction.

✿ implemented by writing transactions into a transaction log



Transaction Processing Services

- ✿ Concurrency control
- ✿ Recovery management
- ✿ Service characteristics

✿ Transparent

- Application programmer does not write code except SQL statements

✿ Consume significant resources

✿ Significant cost component

✿ Transaction design important



Concurrency Control

- ✿ Problem definition
- ✿ Concurrency control problems
- ✿ Concurrency control tools



Concurrency Control Problem

✿ Objective:

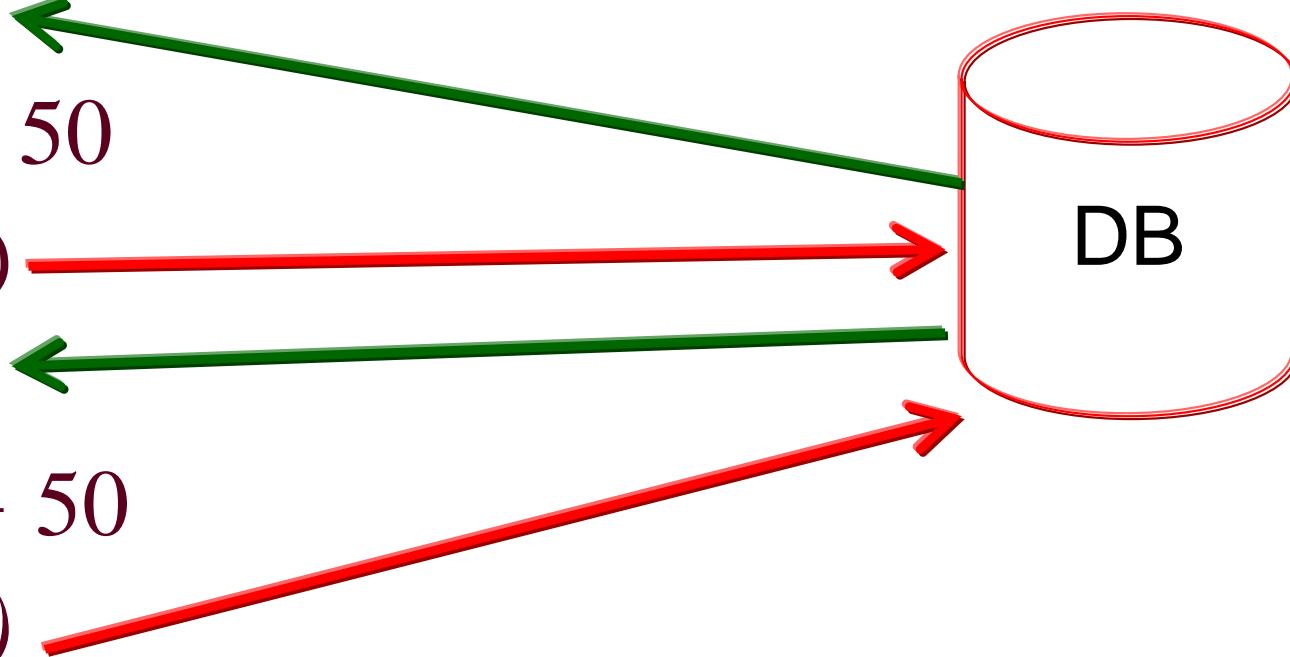
- ✿ Maximize the amount of work performed
- ✿ Throughput: number of transactions processed per unit time(min)
 - Large organizations: thousands of transactions per minute
- ✿ Generating schedules with the Serializability property.
 - **Serializability** of a schedule means equivalence (in the resulting database values) to some *serial* schedule with the same transactions



Example of Fund Transfer

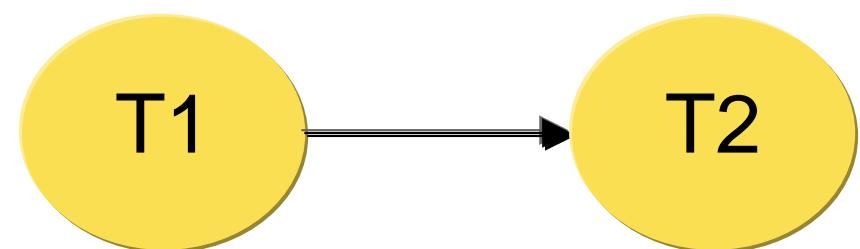
Transaction to transfer \$50 from account A to account B :

1. $\text{read}(A)$
2. $A := A - 50$
3. $\text{write}(A)$
4. $\text{read}(B)$
5. $B := B + 50$
6. $\text{write}(B)$



Serial Schedule

T_1	T_2
<code>read(A)</code> $A := A - 50$ <code>write (A)</code> <code>read(B)</code> $B := B + 50$ <code>write(B)</code>	<code>read(A)</code> $temp := A * 0.1$ $A := A - temp$ <code>write(A)</code> <code>read(B)</code> $B := B + temp$ <code>write(B)</code>



Serializability

S1

T ₁	T ₂
<pre>read(A) A := A - 50 write(A)</pre> <pre>read(B) B := B + 50 write(B)</pre>	 <pre>read(A) temp := A * 0.1 A := A - temp write(A)</pre> <pre>read(B) B := B + temp write(B)</pre>

S2

T ₁	T ₂
<pre>read(A) A := A - 50 write(A)</pre> <pre>read(B) B := B + 50 write(B)</pre>	 <pre>read(A) temp := A * 0.1 A := A - temp write(A)</pre> <pre>read(B) B := B + temp write(B)</pre>

Effects are the same



Concurrency Control Problem

✿ Constraint:

✿ No interference:

- effect is the same as **serial effect** (1-by-1 execution)

✿ Interference occurs on commonly manipulated data known as hot spots

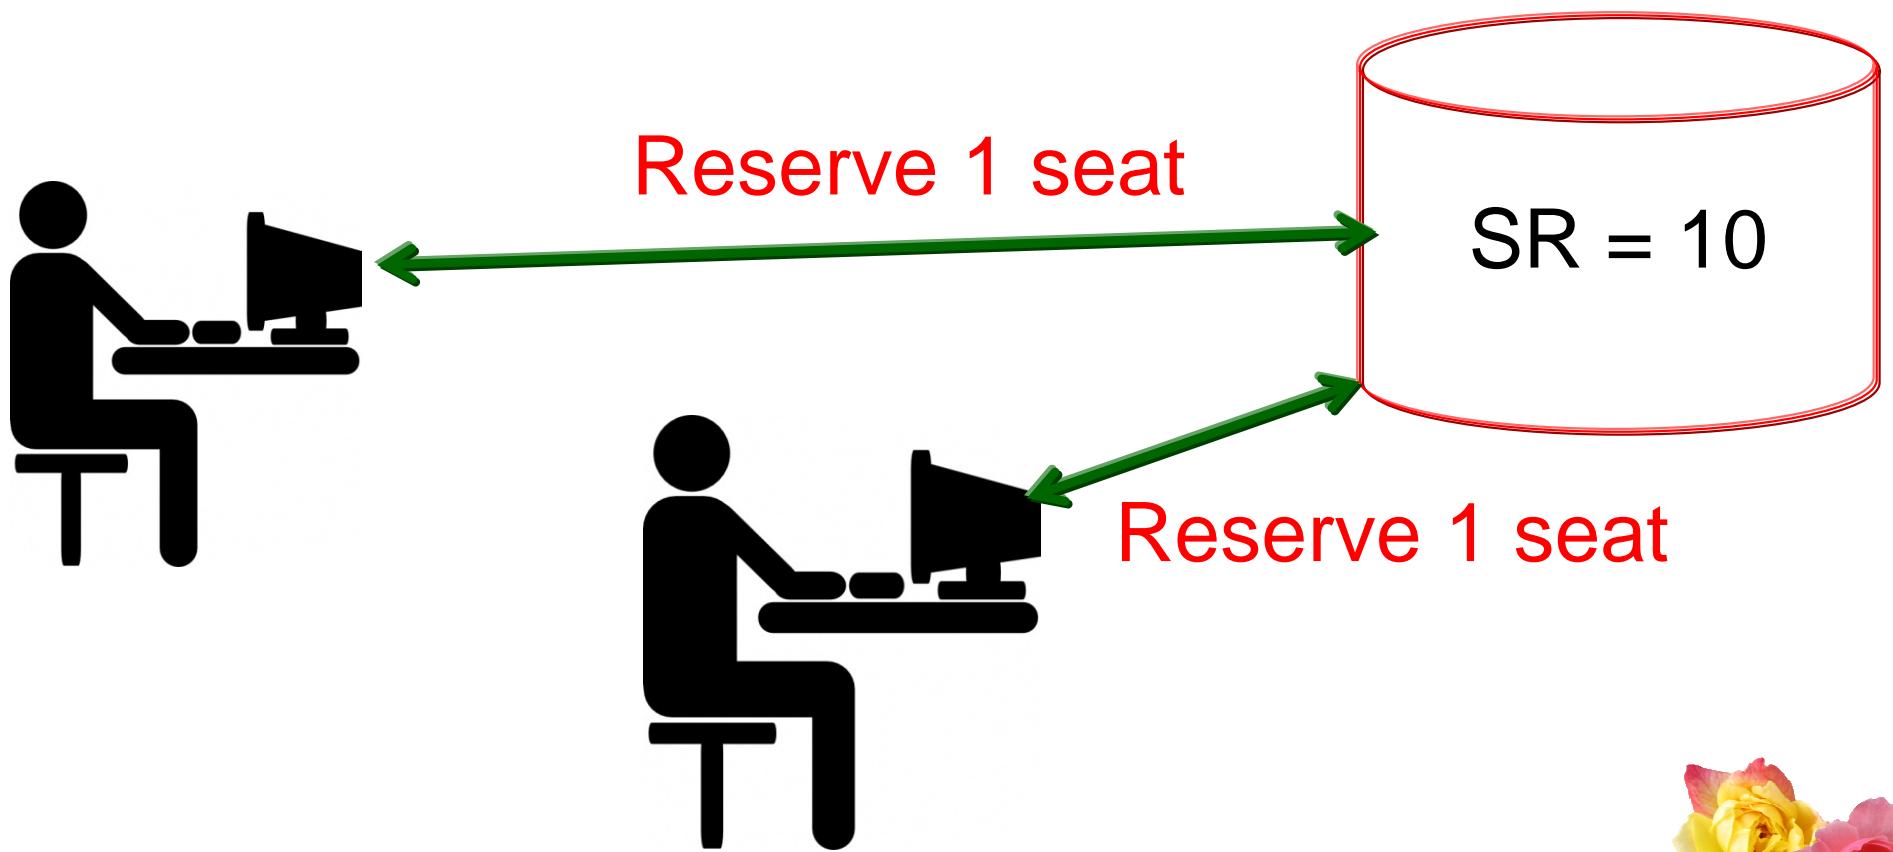
- Example: seats remaining in a flight record

✿ concurrency control must be applied to the entire database, not just the hot spots



Lost Update Problem

Two users trying to **reserve the same flight**:
need to change seats remaining (SR) field



Lost Update Problem

Two users changing the same part of the database

Transaction A	Time	Transaction B
Read SR (10)	T_1	
	T_2	Read SR (10)
If $SR > 0$ then $SR = SR - 1$	T_3	
	T_4	If $SR > 0$ then $SR = SR - 1$
Write SR (9)	T_5	
	T_6	Write SR (9)

SR: Seat Reserved column name

Write-Write problem



Uncommitted Dependency Problem

Transaction A	Time	Transaction B
Read SR (10)	T ₁	
SR = SR - 1	T ₂	
Write SR (9)	T ₃	
	T ₄	Read SR (9)
ROLLBACK ★	T ₅	

failure

the write-read problem



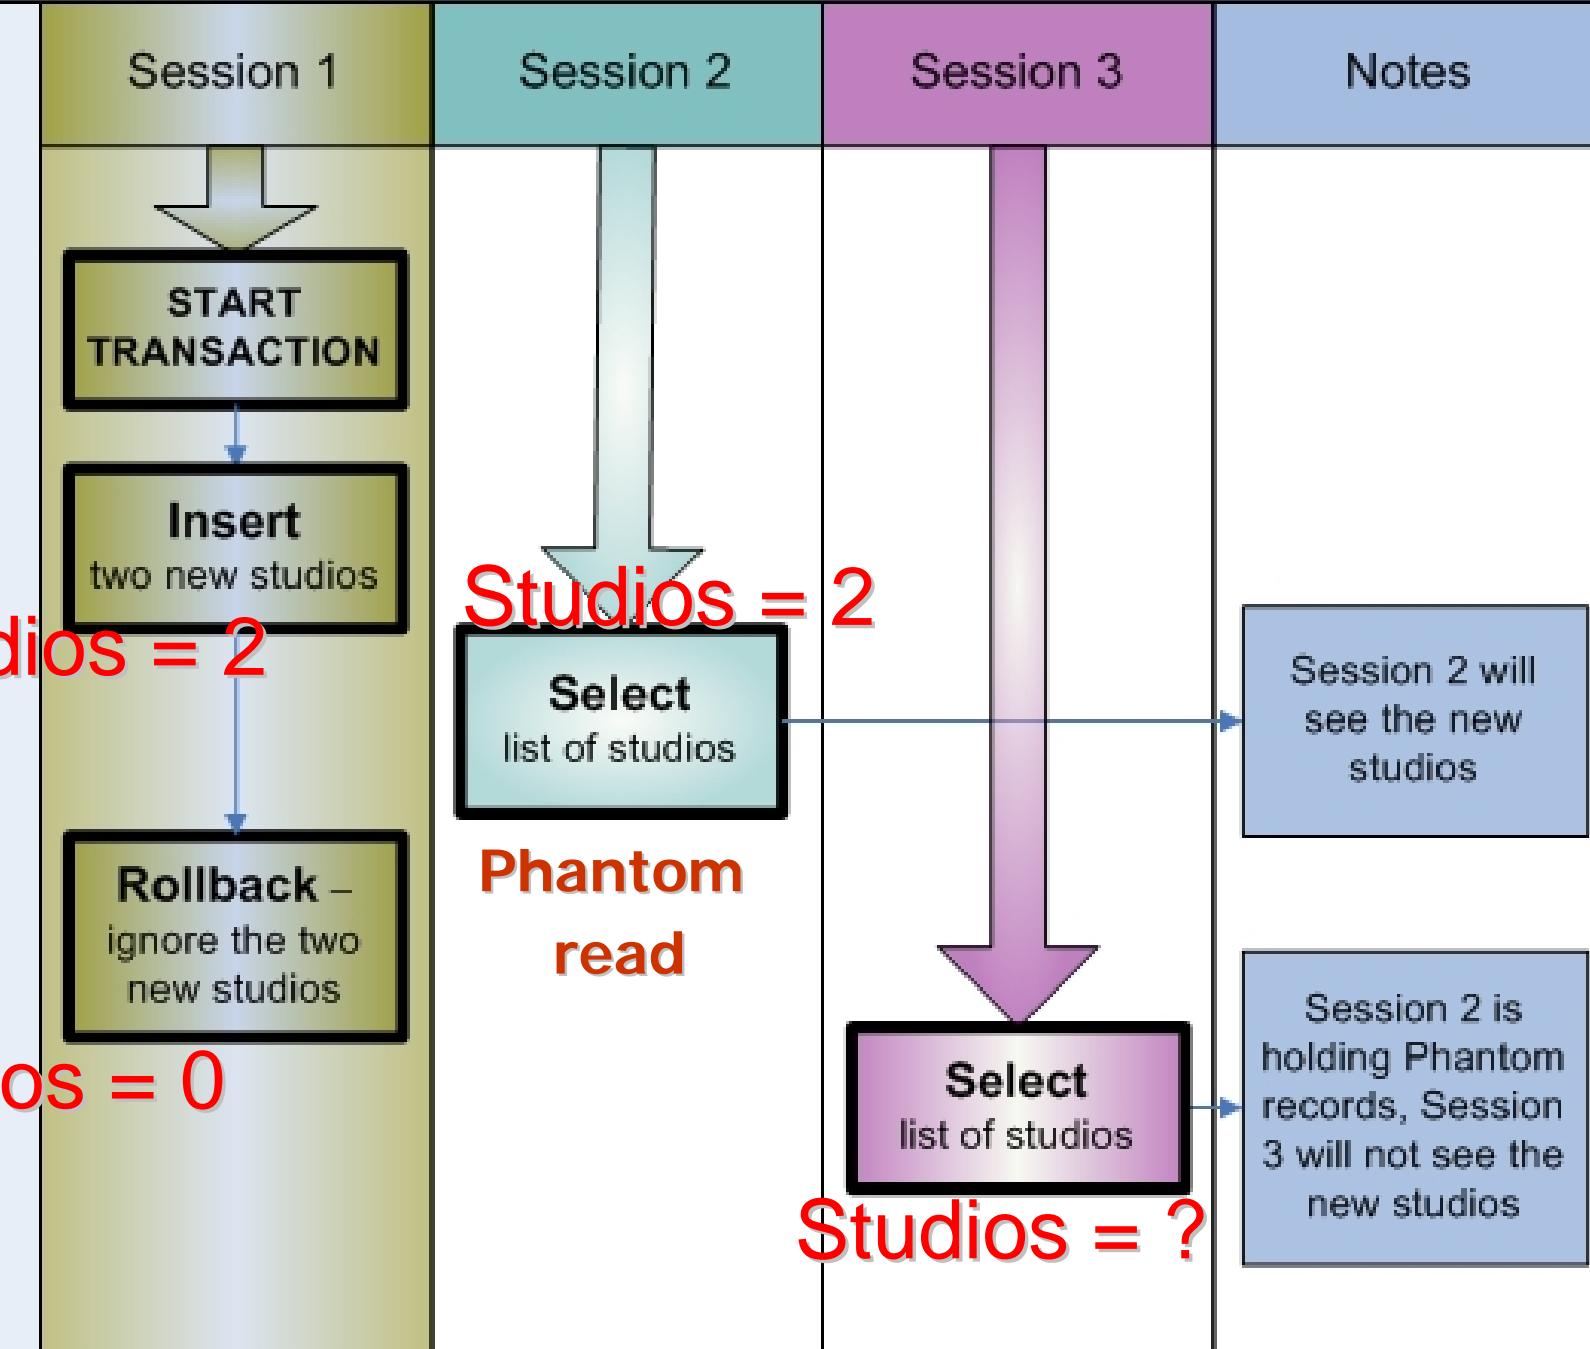
Inconsistent Retrieval Problems

- ✿ Interference causes inconsistency among multiple retrievals of a subset of data
- ✿ Incorrect summary
- ✿ Phantom read
- ✿ Non repeatable read

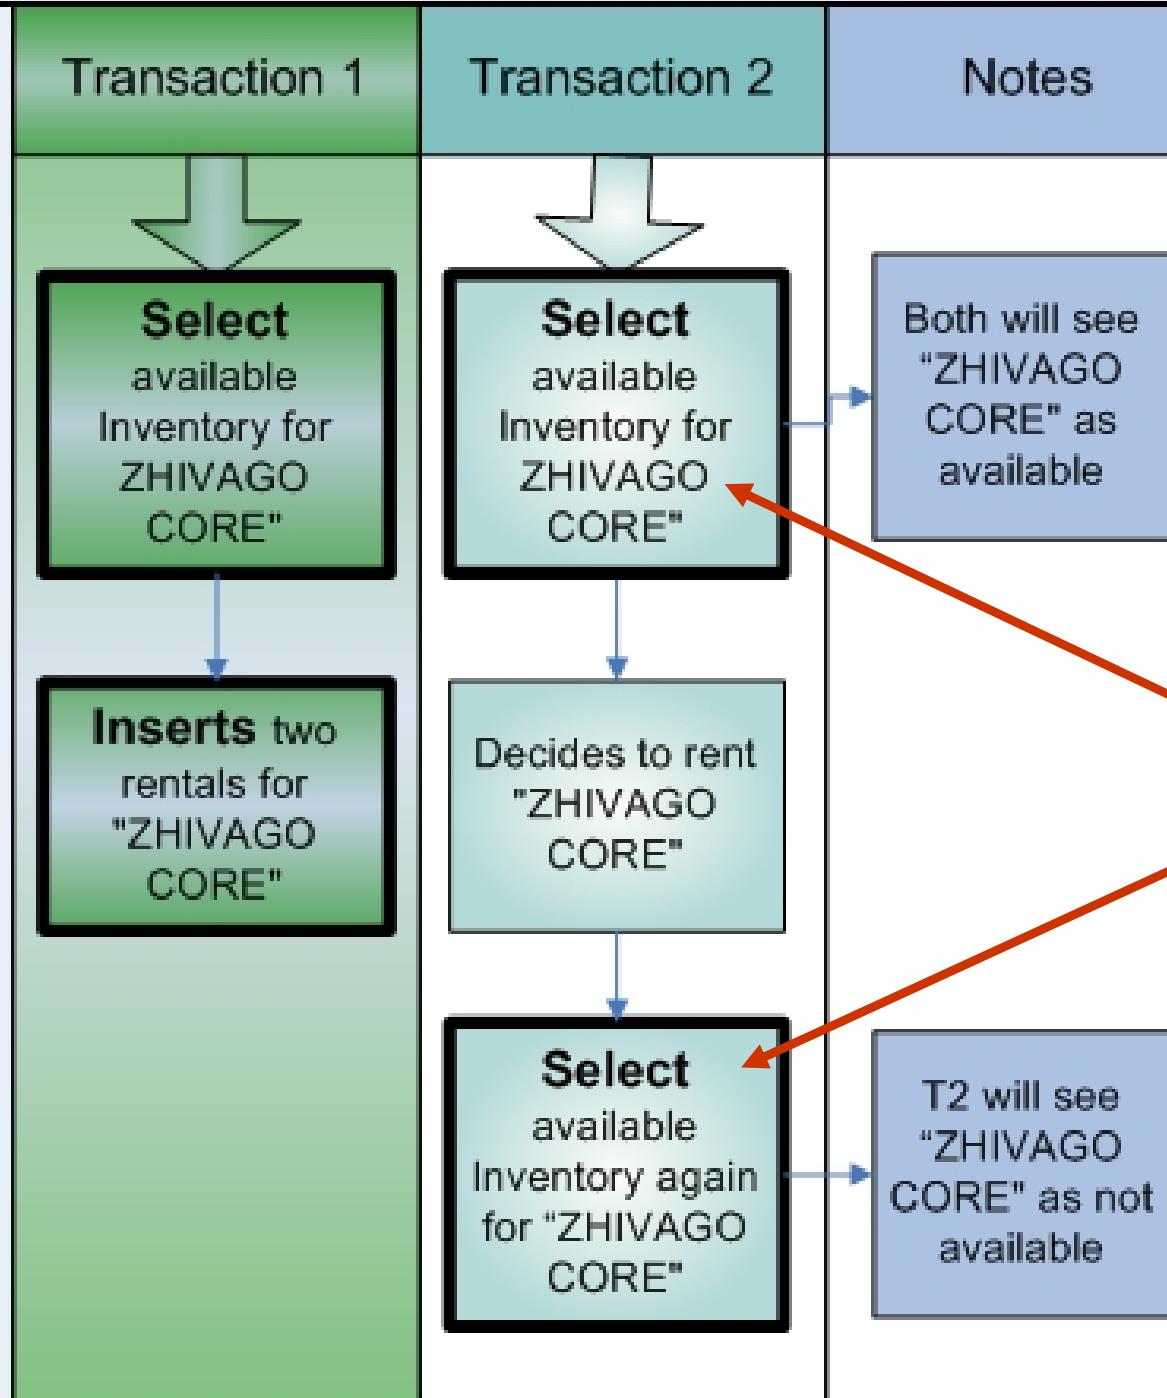


Phantom Reads

Studios = 0

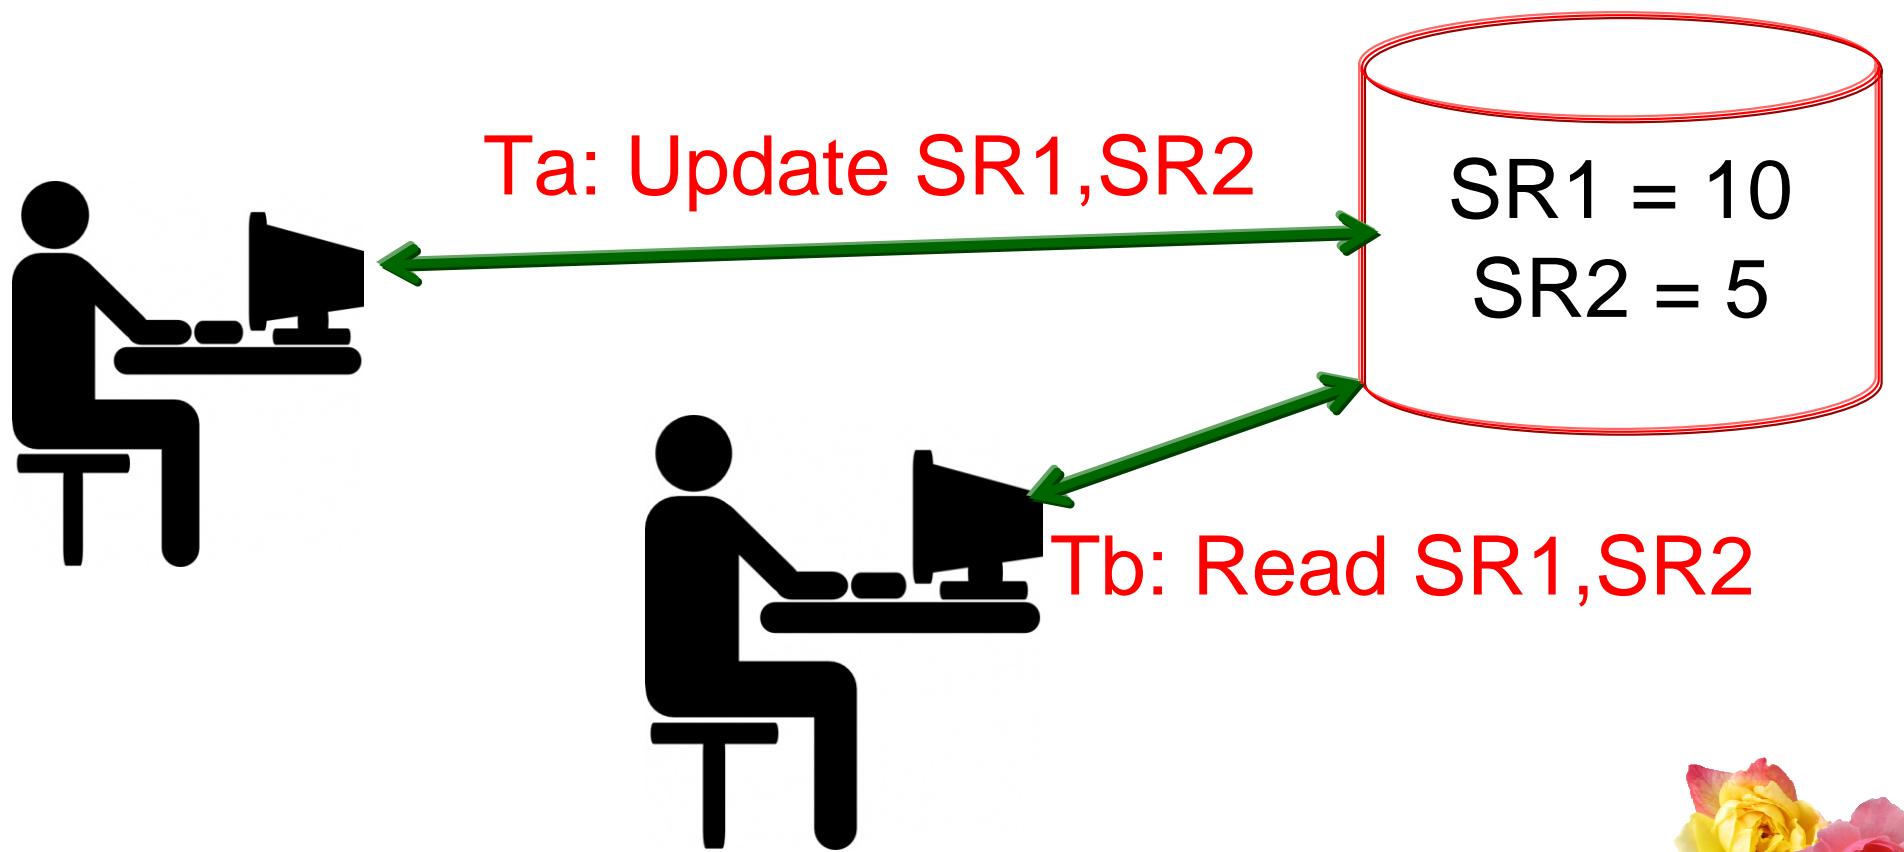


Non-Repeatable Reads

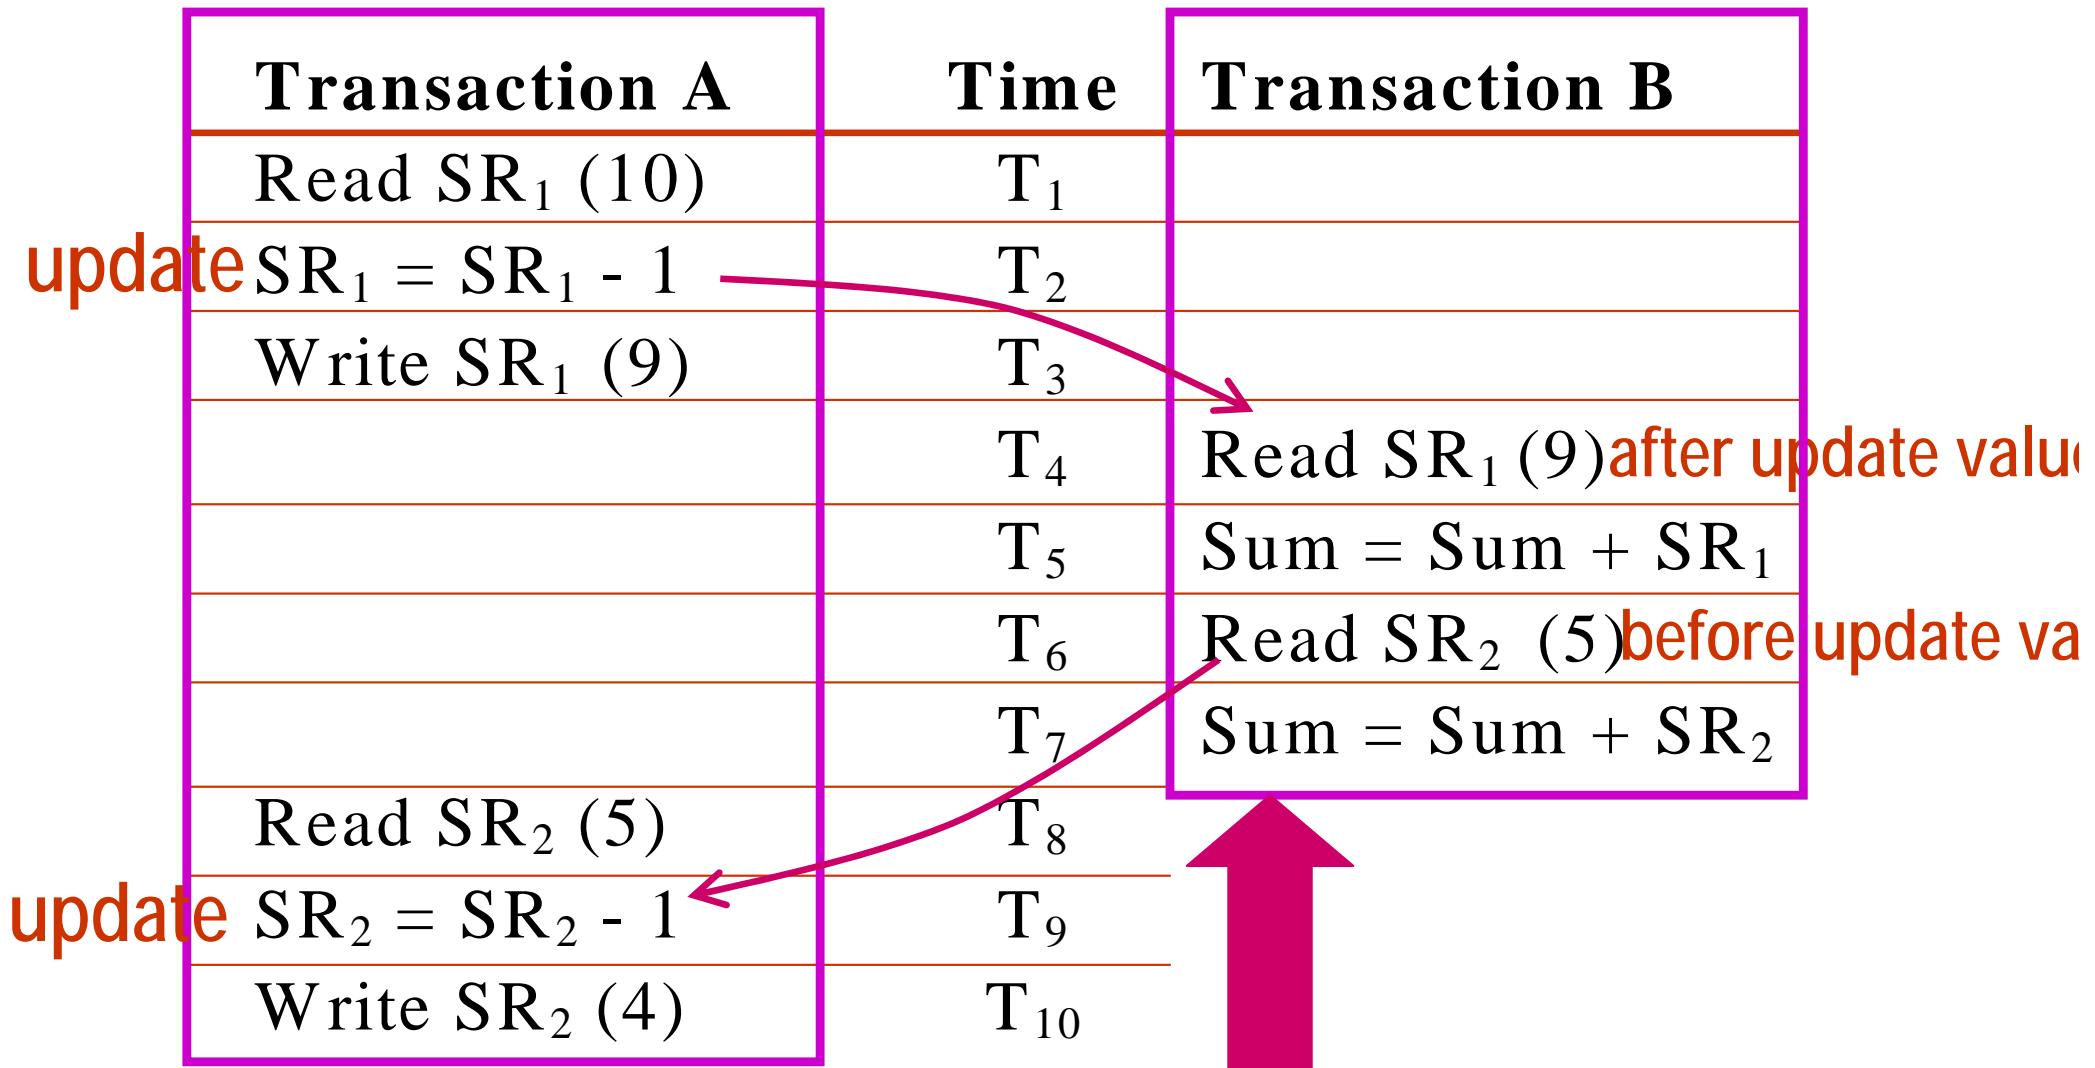


Incorrect Summary Problem

Interference causes inconsistency



Incorrect Summary Problem



Some data in the summary calculation reflects before update values, some data reflects after update values

HOW TO solve concurrency problems ?



Locking Fundamentals

- ✿ Fundamental tool of concurrency control
 - ✿ Locks can be used to prevent interference problems
- ✿ Obtain lock before accessing an item
 - ✿ release lock when finish
- ✿ Wait if a conflicting lock is held
- ✿ Concurrency control manager maintains the lock table

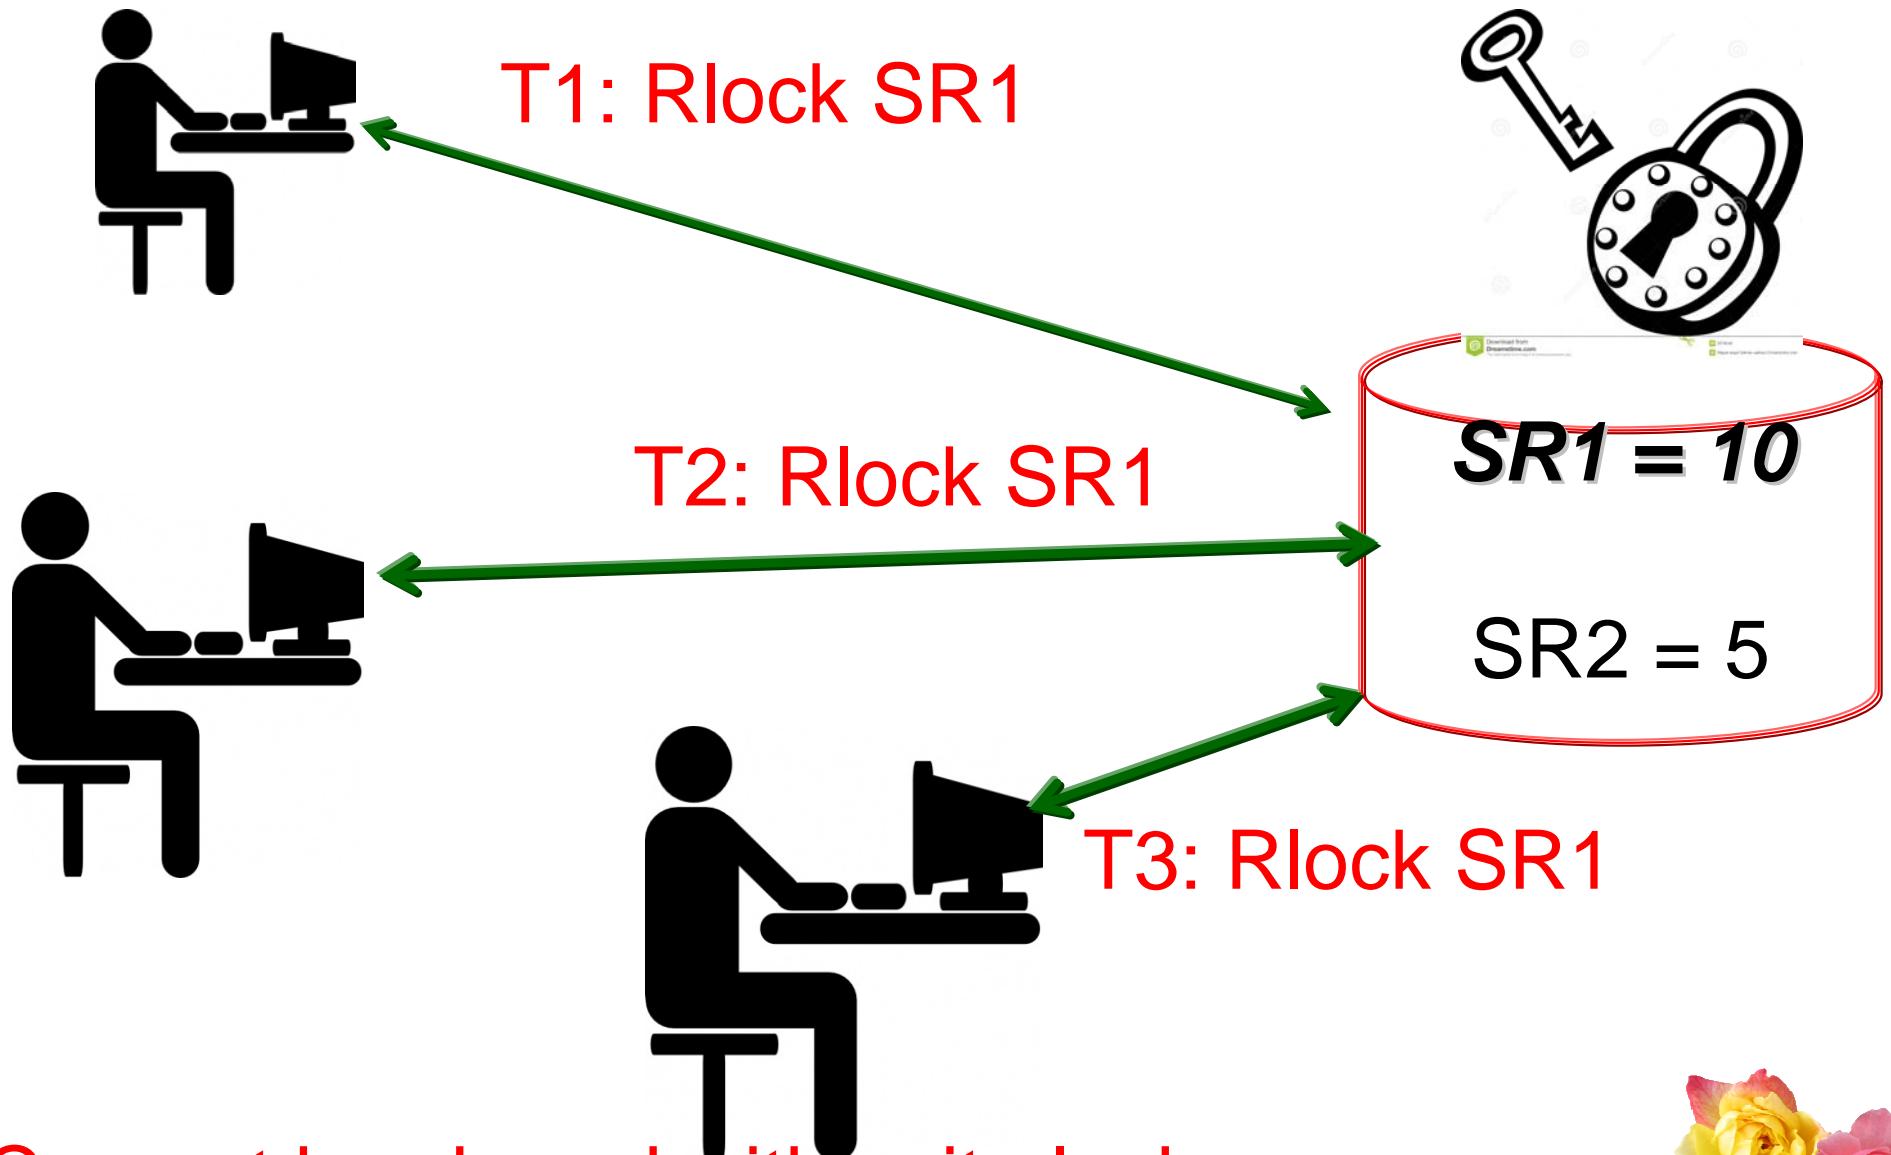


Locking data

- ✿ Shared lock: a read lock
- ✿ Exclusive lock: a write lock
- ✿ Not allow to share Xlock



Shared lock: a read lock



Can not be shared with write lock



Locking Fundamentals

✿ Conflicting locks

✿ **Shared lock:** conflicts with exclusive locks:

- Write lock

✿ **Exclusive lock:** conflicts with all other kinds of locks:

- Write lock
- Read lock



Locking Fundamentals

✿ Lock usage:

- ✿ Obtain appropriate kind of lock before accessing database item
- ✿ Wait if another transaction holds a conflicting lock

✿ Lock table:

- ✿ Details about locks held by transactions
- ✿ Concurrency control manager maintains
- ✿ Lock: insert a lock record
- ✿ Release (unlock): delete a lock record

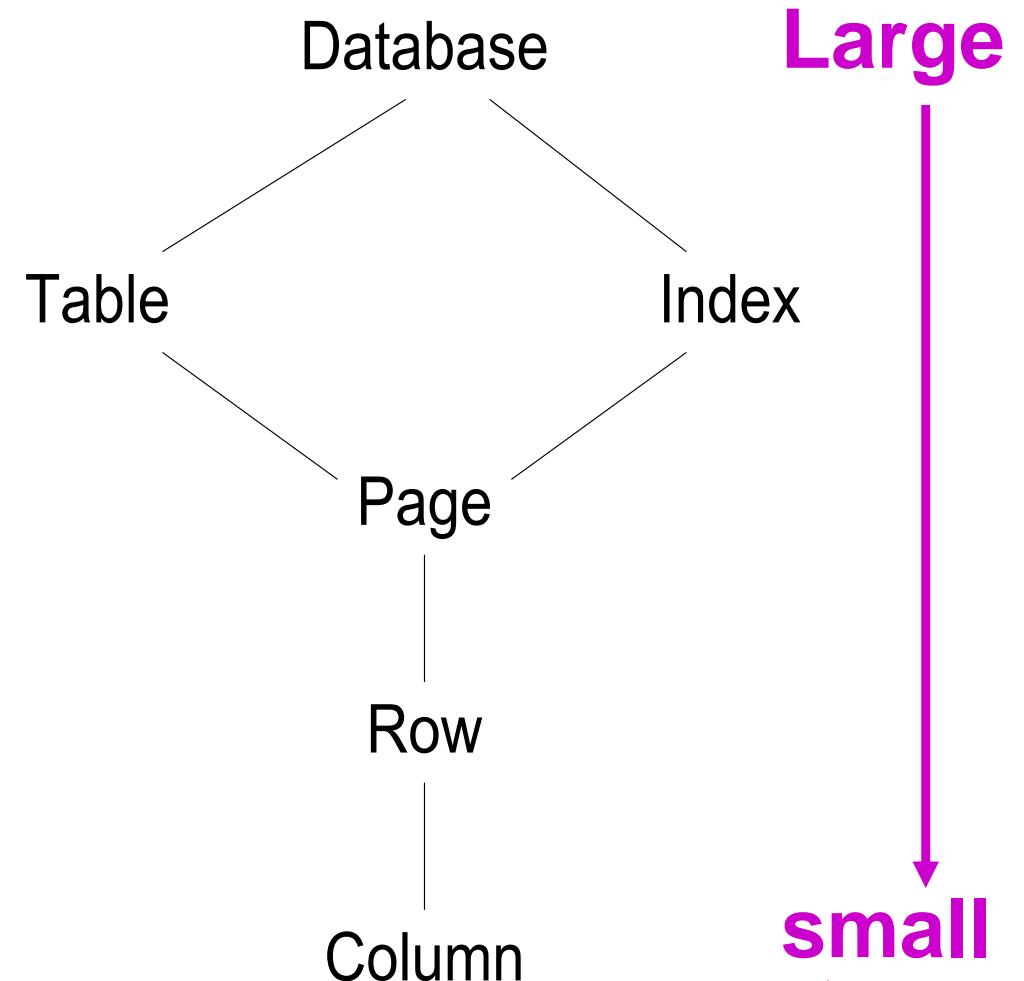


Locking Granularity

(size of database item locked)

✿ Tradeoff:

✿ Overhead (#of locks)
versus waiting



Locking Granularity?

✿ Intent locks (intent-for-update locks)

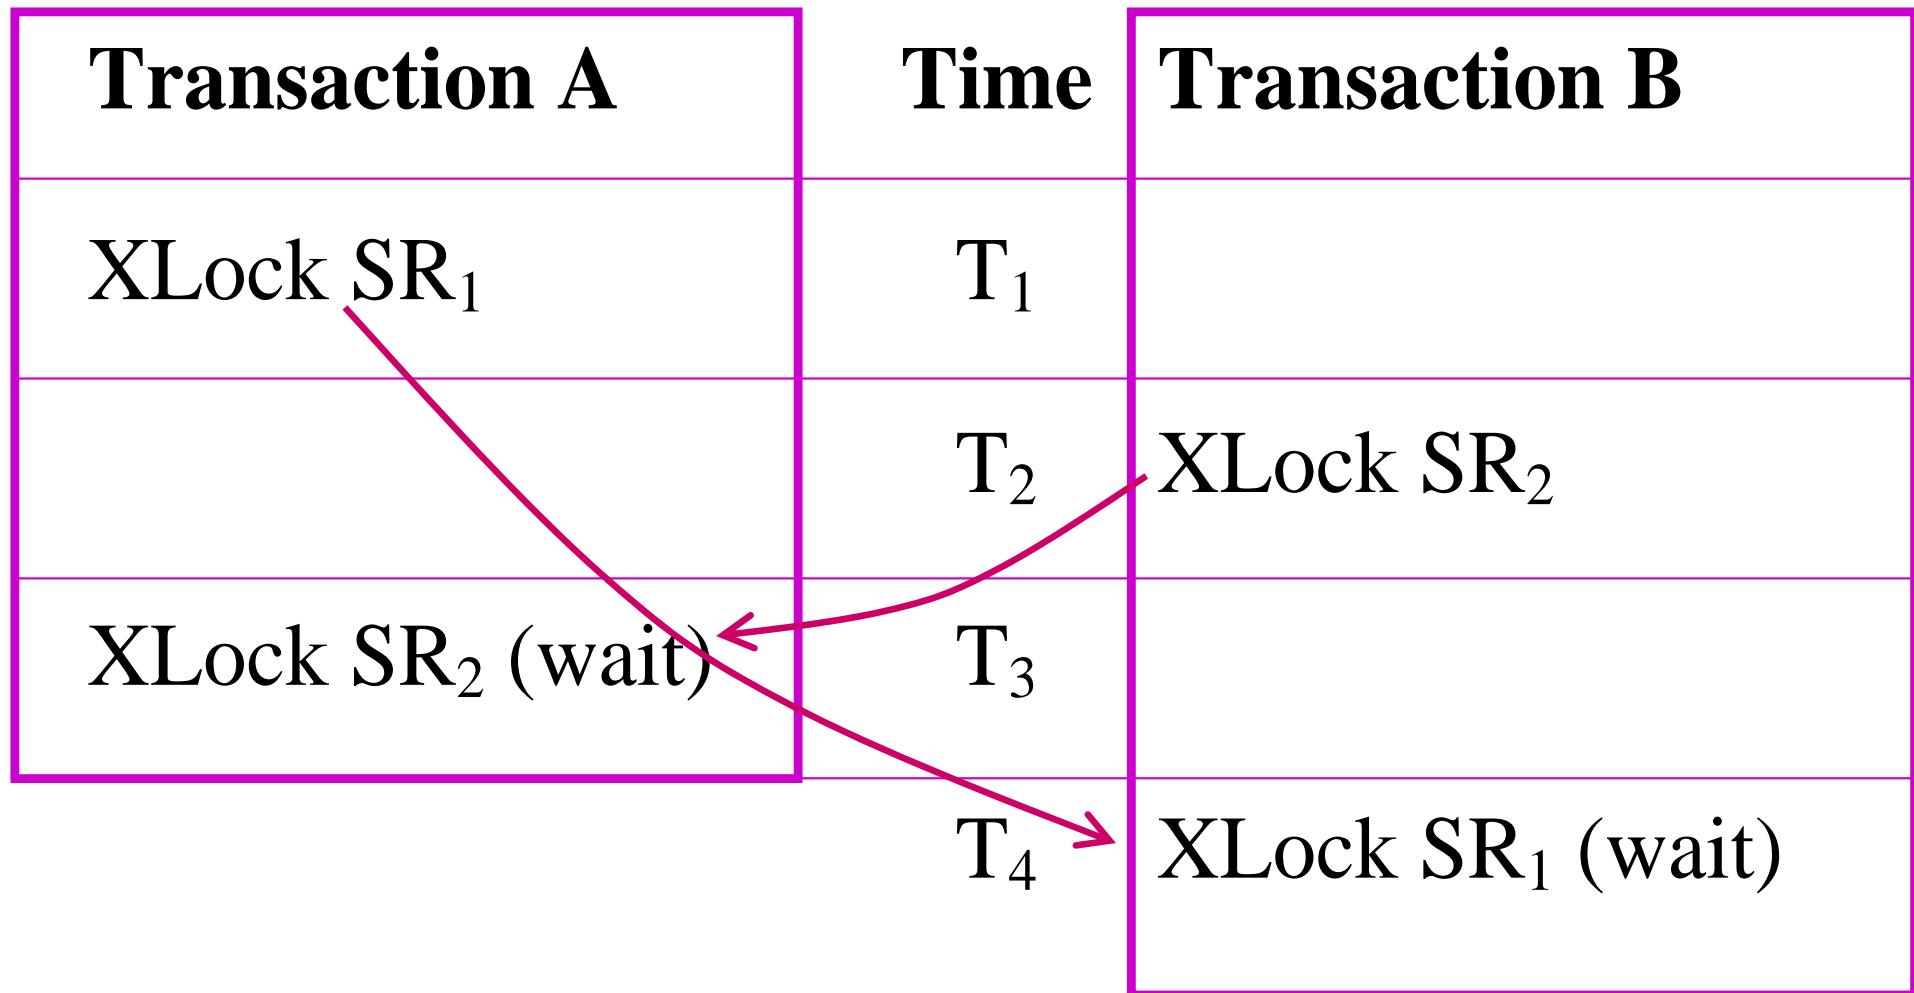
✿ Permit more concurrency control on large items (e.g. table)

✿ Purpose

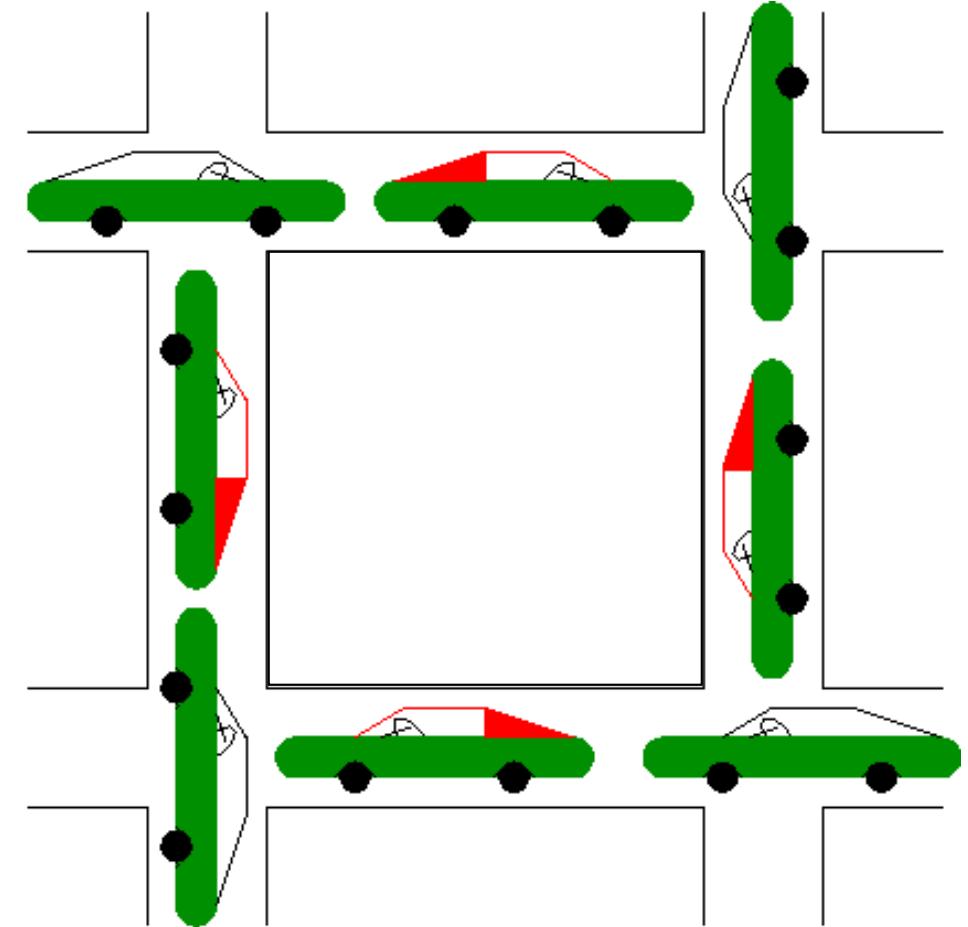
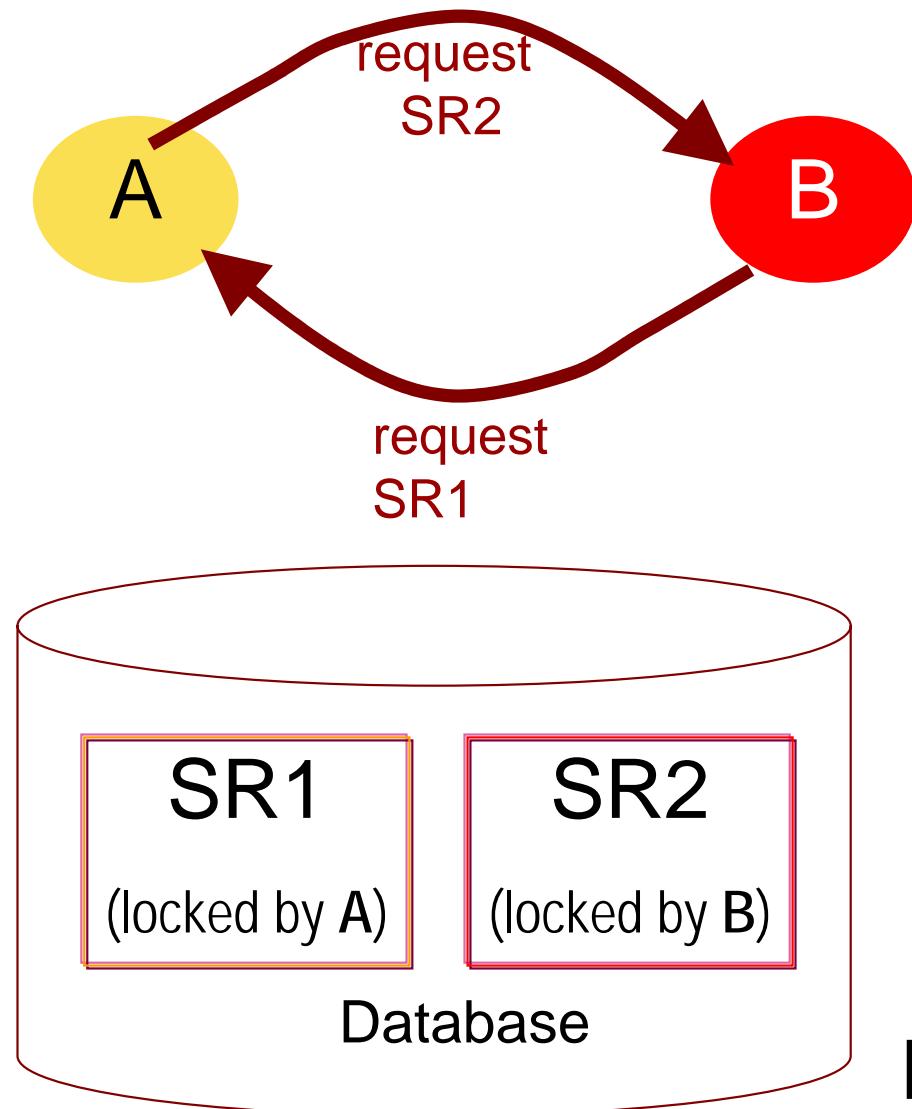
- To prevent other transactions from modifying the higher-level resource in a way that would invalidate the lock at the lower level.
 - Obtain intent lock on large item before shared/exclusive locks on smaller items
- To improve the efficiency of the Database Engine in detecting lock conflicts at the higher level of granularity.



Deadlock (Mutual Waiting)



Deadlock (Mutual Waiting)



Both Ta,Tb want to access
SR1,SR2



Deadlock Resolution

✿ Detection

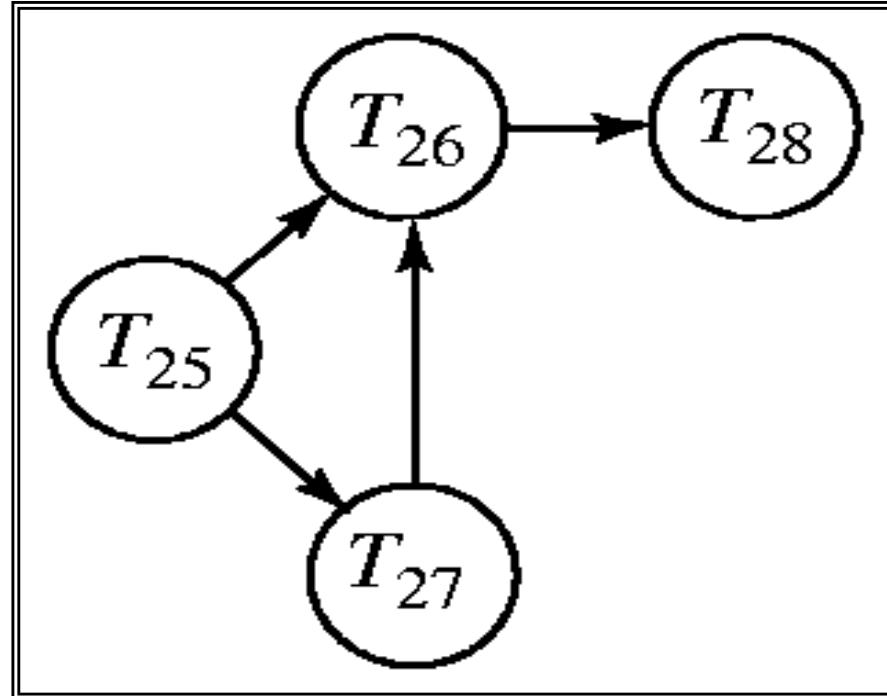
- ✿ Create a graph showing waiting dependencies
 - Search the graph for cycles
- ✿ Empirically, most deadlocks involve 2 or 3 transactions
- ✿ Used by enterprise DBMSs

✿ Prevention: Timeout

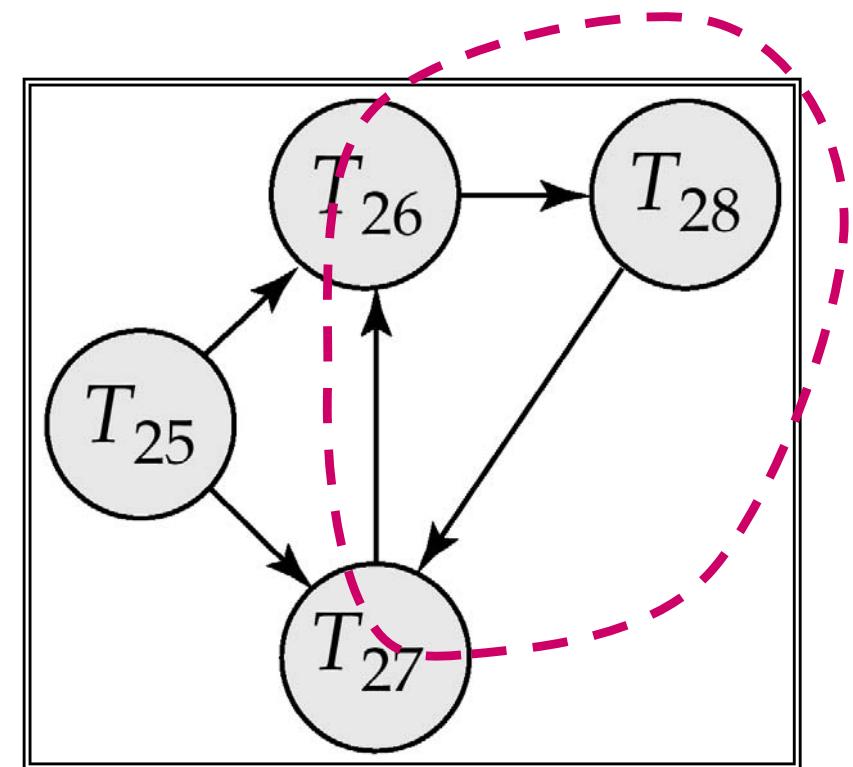
- ✿ Waiting limit
- ✿ Can abort transactions that are not deadlocked
- ✿ Timeout interval is difficult to determine
- ✿ Used by desktop DBMSs



Deadlock Detection (Cont.)



Wait-for graph without a cycle



Wait-for graph with a cycle



Two Phase Locking (2PL)

- ✿ Protocol to prevent *lost update problems*
 1. Locking phase
 2. Unlocking phase

Lock A

....

Lock B

Lock C

....

Unlock B

Unlock A

Unlock C



Two Phase Locking (2PL)

- ✿ Protocol to prevent *lost update problems*
- ✿ All transactions must follow conditions
 - ✿ Obtain lock before accessing item
 - ✿ Wait if a conflicting lock is held
 - ✿ Cannot obtain new locks after releasing locks (within a transaction)



Two distinct consecutive phases

during the transaction's execution:

✿ **Expanding phase** (Growing phase):

 locks are acquired and no locks are released
 (the number of locks can only increase).

✿ **Shrinking phase**:

 locks are released and no locks are acquired.

Guarantee Serializability



2PL

Lock..

...

Lock..

•

•

•

Unlock..

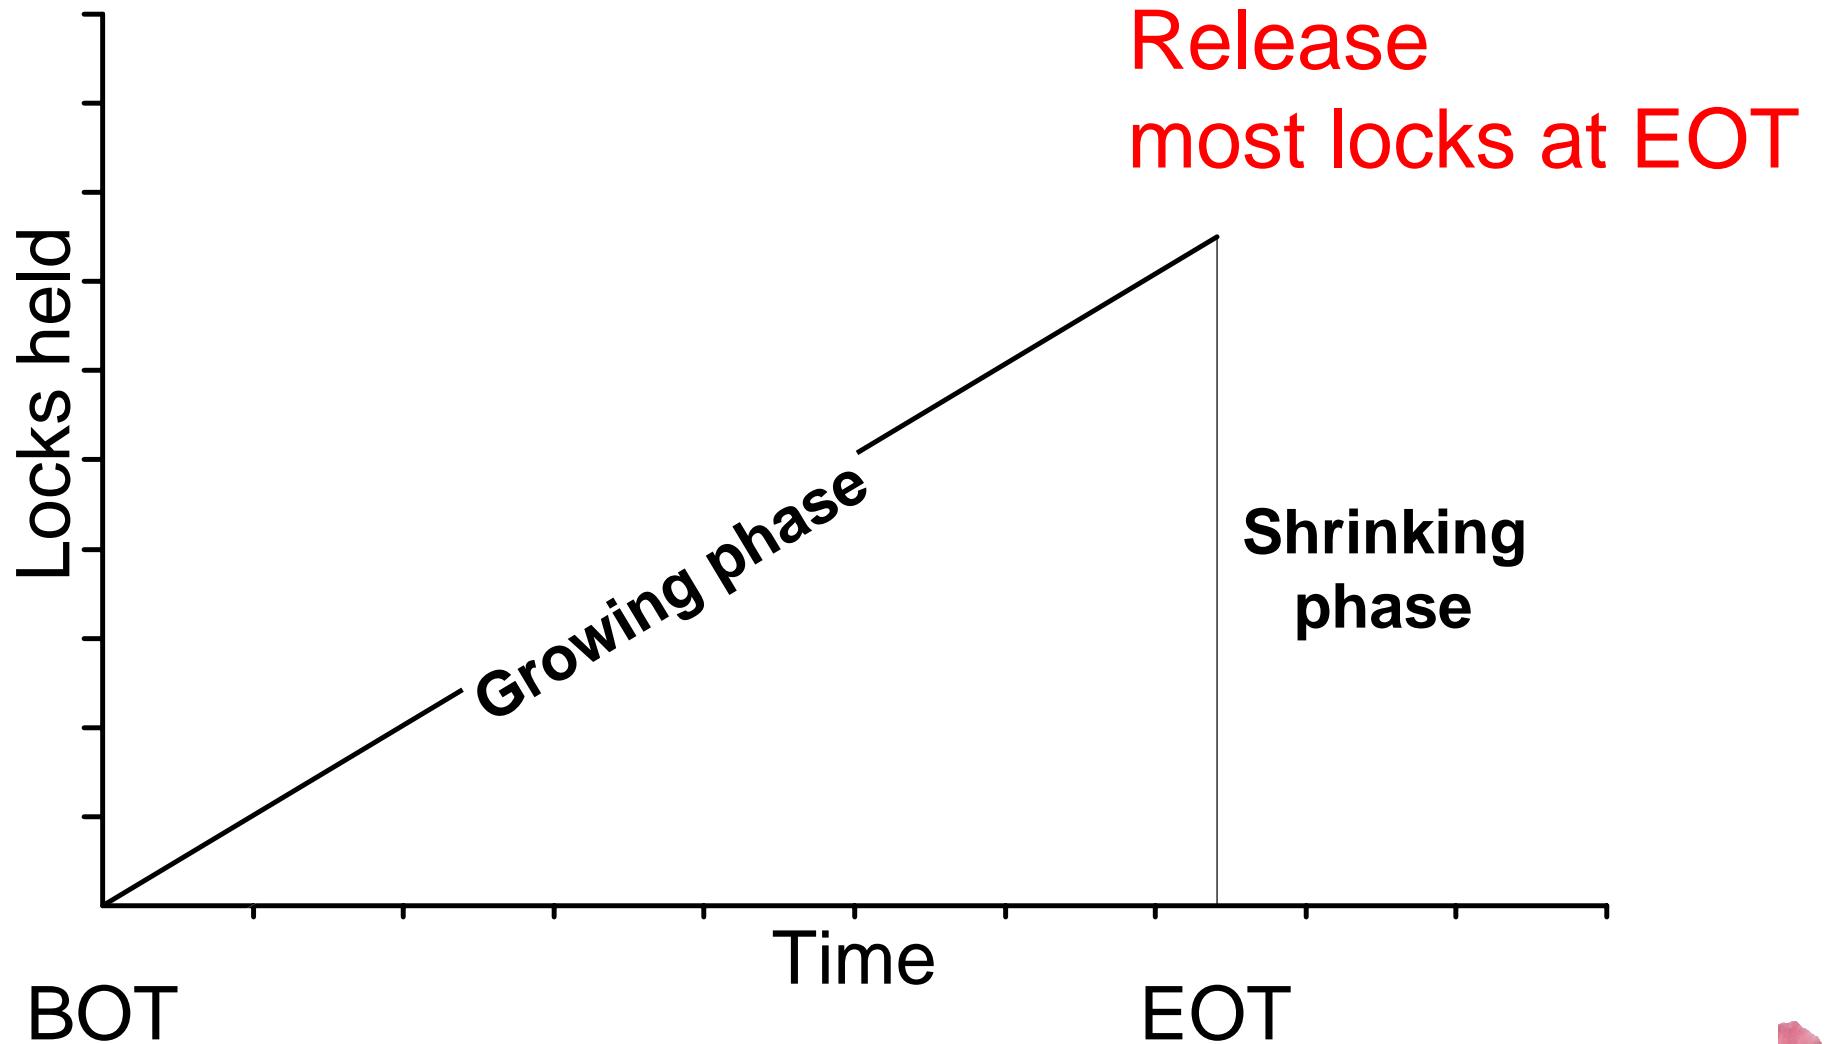
...

Unlock..

T1	T2
lock(A)	
	read(A)
	lock(A)
write(A)	
lock(B)	
unlock(A)	
	read(A)
	write(A)
	lock(B)
	unlock(A)
read(B)	
write(B)	
unlock(B)	
	read(B)
	write(B)
	unlock(B)



2PL Implementation



Optimistic Approaches

- ✿ Assumes conflicts are rare
- ✿ No locks
- ✿ Check for conflicts (version/time control)
 - ✿ After each read and write
 - ✿ At the end of transaction
- ✿ Evaluation
 - ✿ Less overhead: *no locks, deadlocks, ...*
 - ✿ More variability: *penalty for conflict is rollback*
 - ✿ Locking penalty: *waiting (less penalty than rollback)*



Timestamp-Based Protocols

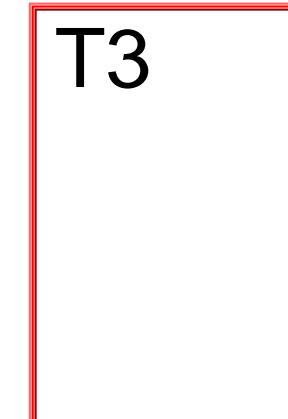
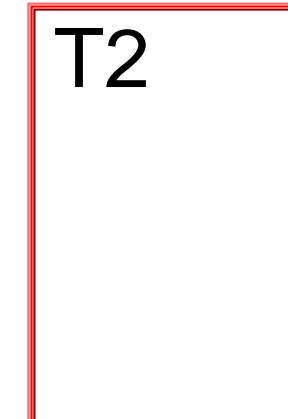
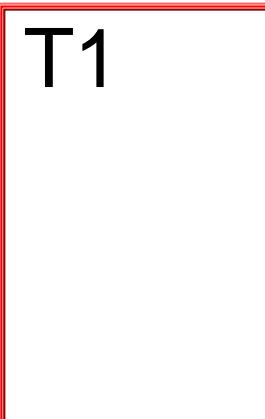
- ✿ Each transaction is issued a timestamp when it enters the system. ($T_{\text{before}} < T_{\text{after}}$)
- ✿ time-stamps determine the serializability order.

Timestamp

101

102

103



Timestamp-Based Protocols

✿ to assure **serializability**, the protocol maintains **for each data Q two timestamp values:**

✿ **W-timestamp(Q)**

- largest time-stamp of any T that executed $\text{write}(Q)$ successfully.

✿ **R-timestamp(Q)**

- largest time-stamp of any T that executed $\text{read}(Q)$ successfully.



101

T1

W(SR1)

102

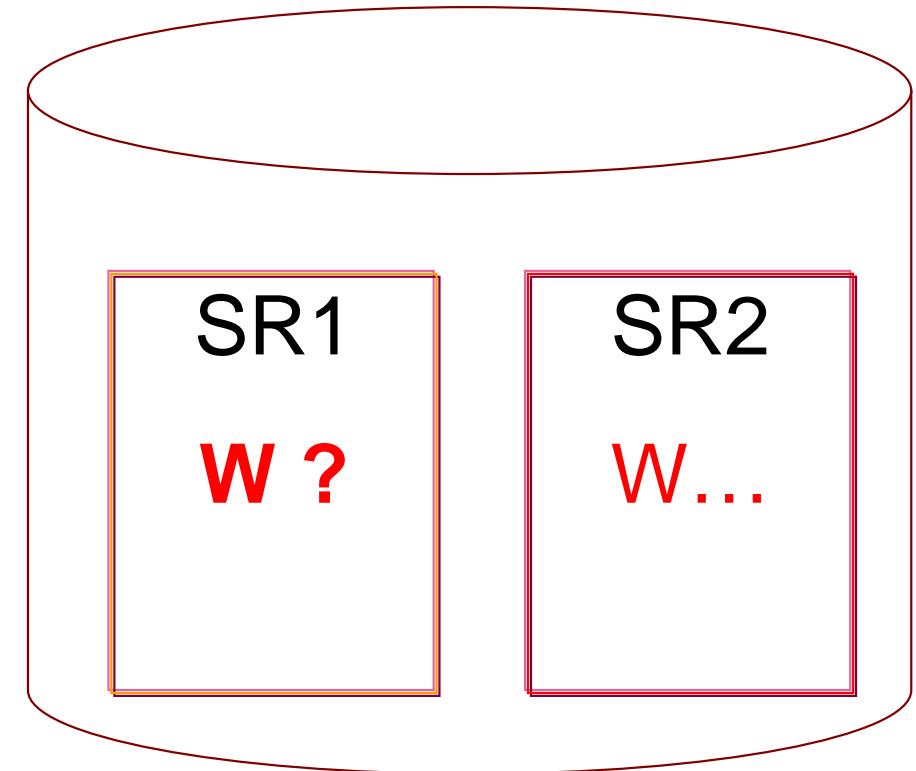
T2

W(SR1)

103

T3

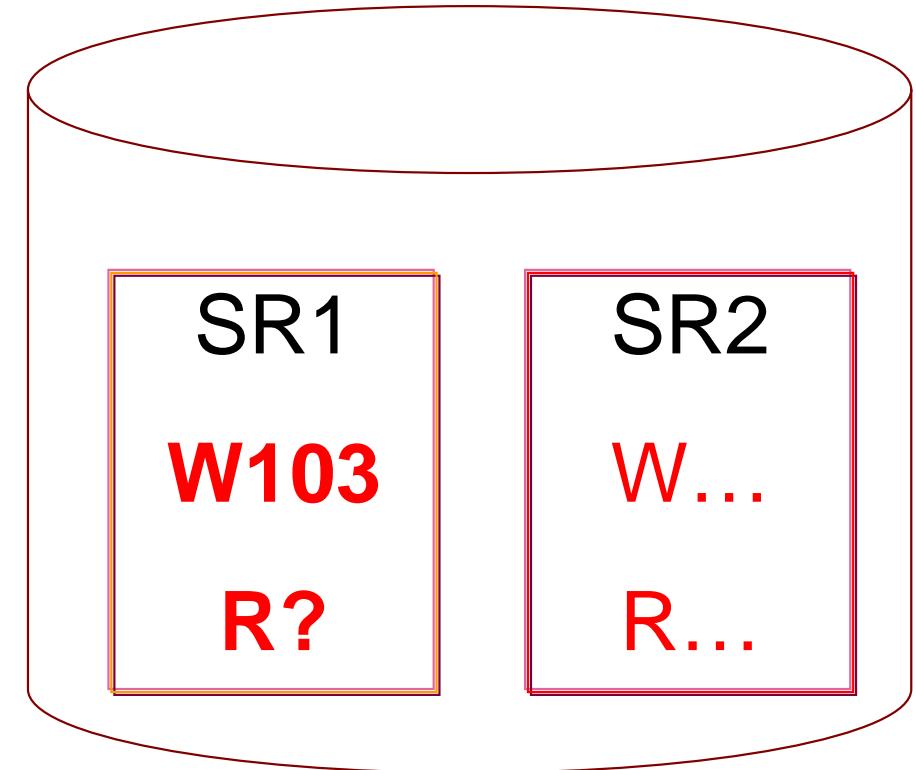
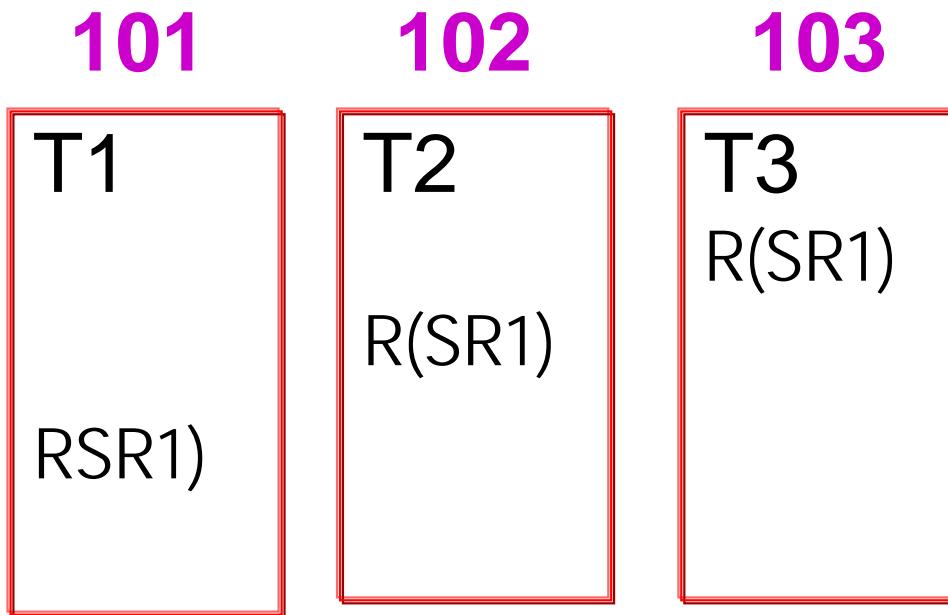
W(SR1)



W-timestamp(Q)

largest time-stamp of any T that executed write(Q) successfully.





R-timestamp(Q)

largest time-stamp of any T that executed $\text{read}(Q)$ successfully.



Timestamp

read/write in the order of timestamps

- ✿ T1, read data(A)
 - ✿ If T2 written A, abort transaction T1-rollback
- ✿ T1, write data(A)
 - ✿ If T2 read A or
 - ✿ T2 written A
 - Abort transaction T1
 - ✿ Other cases, T1 can write A
- ✿ After read/write, update timestamp A



Timestamp rules

- ✿ Maintain the serial order
 - ✿ T1 must write data before T2 read, write
 - ✿ T1 can not read data that written by T2.

Write, Write
Read, Write



Example : *Timestamp approach*

A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5

T_1	T_2	T_3	T_4	T_5	
read(Y)	read(Y)			read(X)	Xw , Xr
	read(X)	write(Y) write(Z)		read(Z)	Yw , Yr Zw , Zr
read(X)		write(Z) abort		write(Y) write(Z)	



Locking strategies by table type.

Table Type	Examples	Suggested Locking Strategy
Live-High Volume	• Account	• <u>Optimistic</u> (first choice) • <u>Pessimistic</u> (second choice)
Live-Low Volume	• Customer • Insurance Policy	• <u>Pessimistic</u> (first choice) • <u>Optimistic</u> (second choice)
Log (<i>typically append only</i>)	• AccessLog • AccountHistory • TransactionRecord	• <u>Overly Optimistic</u>
Lookup/Reference (<i>typically read only</i>)	• State • PaymentType	• <u>Overly Optimistic</u>

Overly Optimistic is suitable for single user systems



Recovery Management

✿ Device characteristics and failure types

✿ Recovery tools

 ✿ *transaction log, checkpointing*

✿ Recovery processes



Storage Device Basics

- ✿ Volatile: loses state after a shutdown (*RAM*)
- ✿ Nonvolatile: retains state after a shutdown (*Disk, tape, optical media*)
- ✿ Nonvolatile is more reliable than volatile but failures can cause loss of data
- ✿ Use multiple levels and redundant levels of nonvolatile storage for valuable data



Failure Types

✿ Local

- ✿ Program detected and abnormal termination(/zero)
- ✿ Affects only a single transaction

✿ Operating System

- ✿ Affects all active transactions
- ✿ Less common than local failures

✿ Device (*disk failure*)

- ✿ Affects all active and past transactions
- ✿ Least common



Transaction Log

- ✿ History of database changes
- ✿ Large storage overhead
- ✿ Operations
 - ✿ Undo: revert to previous state
 - ✿ Redo: reestablish a new state
- ✿ Fundamental tool of recovery management



Transaction Log Example

LSN	TransNo	Action	Time	Table	Row	Column	Old	New
1	101001	START	10:29					
2	101001	UPDATE	10:30	Acct	10001	AcctBal	100	200
3	101001	UPDATE	10:30	Acct	15147	AcctBal	500	400
4	101001	INSERT	10:32	Hist	25045	* <i>all columns</i>	<1002, 500, ...>	
5	101001	COMMIT	10:33					



Checkpoints

- ✿ Reduce restart work but adds overhead
 - ✿ Checkpoint log record
(record working transaction at the checkpoint)
 - ✿ Write log buffers and database buffers (to disk)
- ✿ Checkpoint interval: time between checkpoints (*5 to 10 minutes in high volume environments*)
- ✿ All transaction processing stops during the checkpoint process



Other Recovery Tools

✿ Force writing to disk when

- ✿ Checkpoint time

- ✿ End of transaction

✿ Database backup

- ✿ Complete, entire database

- ✿ Incremental



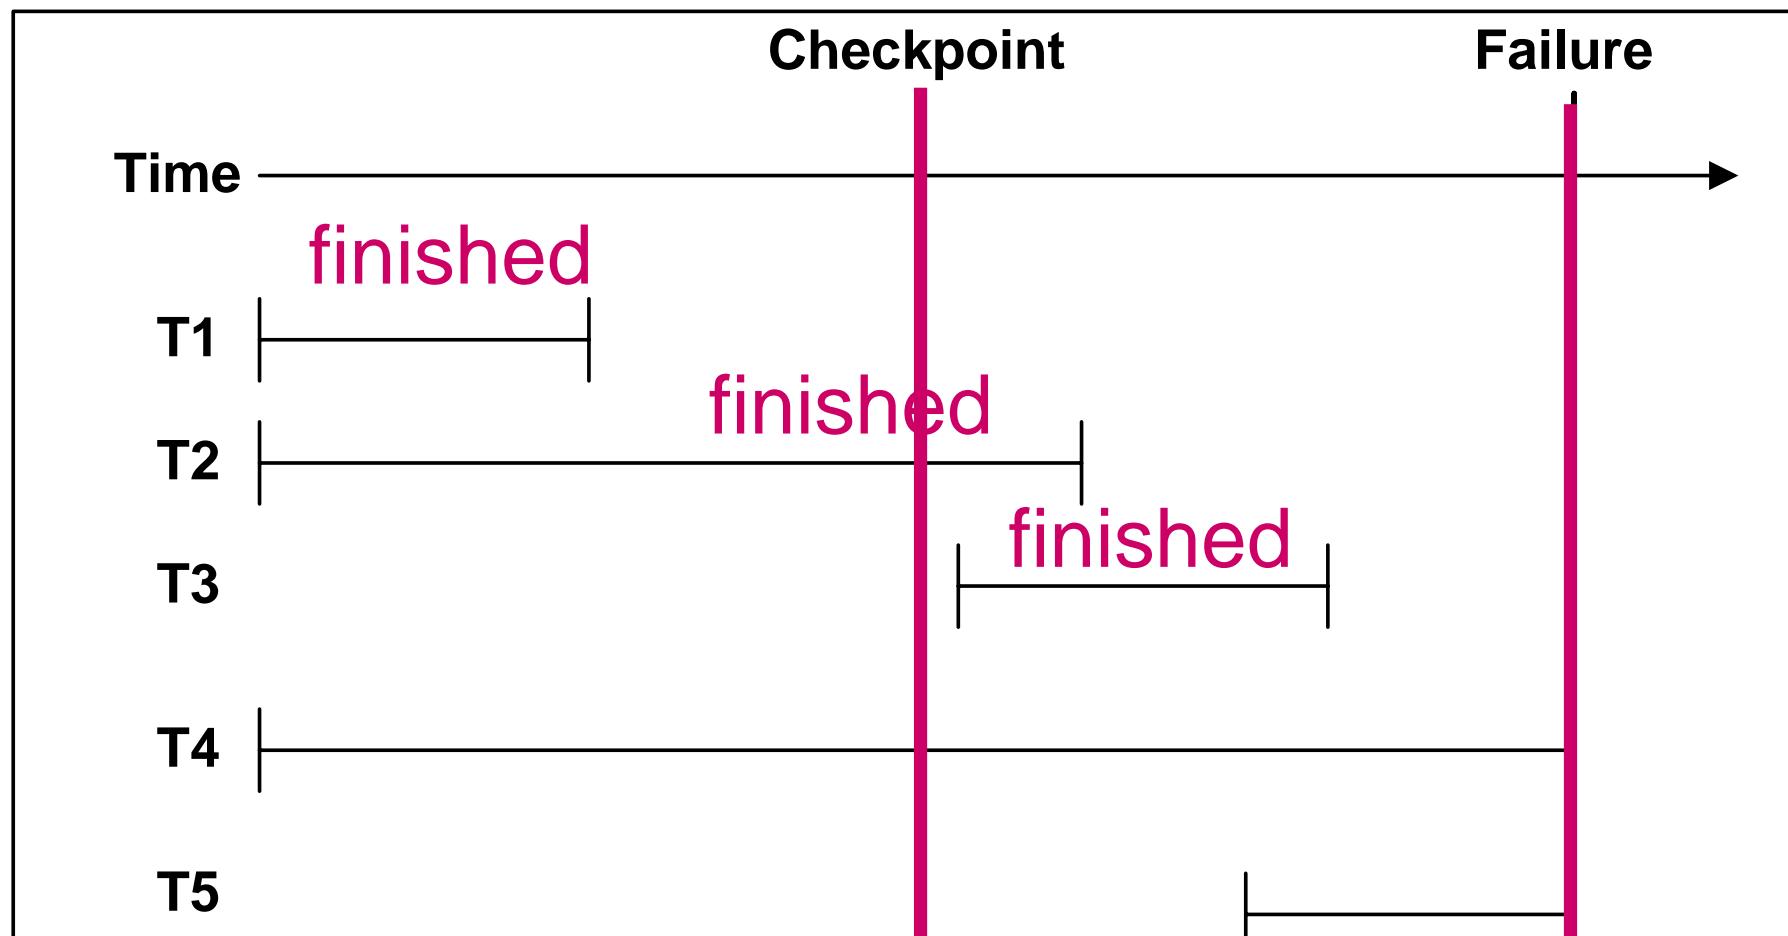
Recovery from a Media Failure

- ✿ Restore database from the most recent backup
- ✿ Redo all committed transactions since the most recent backup
- ✿ Restart active transactions
 - ✿ Rollback and then restart

Assumes not log failure



Recovery Timeline



Recovery Processes

- ✿ Depend on timing of database writes
- ✿ Immediate update approach:
 - ✿ Before commit (*update database before committed*)
 - ✿ Log records written to disk first
 - ✿ Redo, undo
- ✿ Deferred update approach
 - ✿ After commit (*wait for commit before update database*)
 - ✿ Undo operations not needed



Recovery process

✿ Redo process

✿ read log file and redo operations recorded in log file

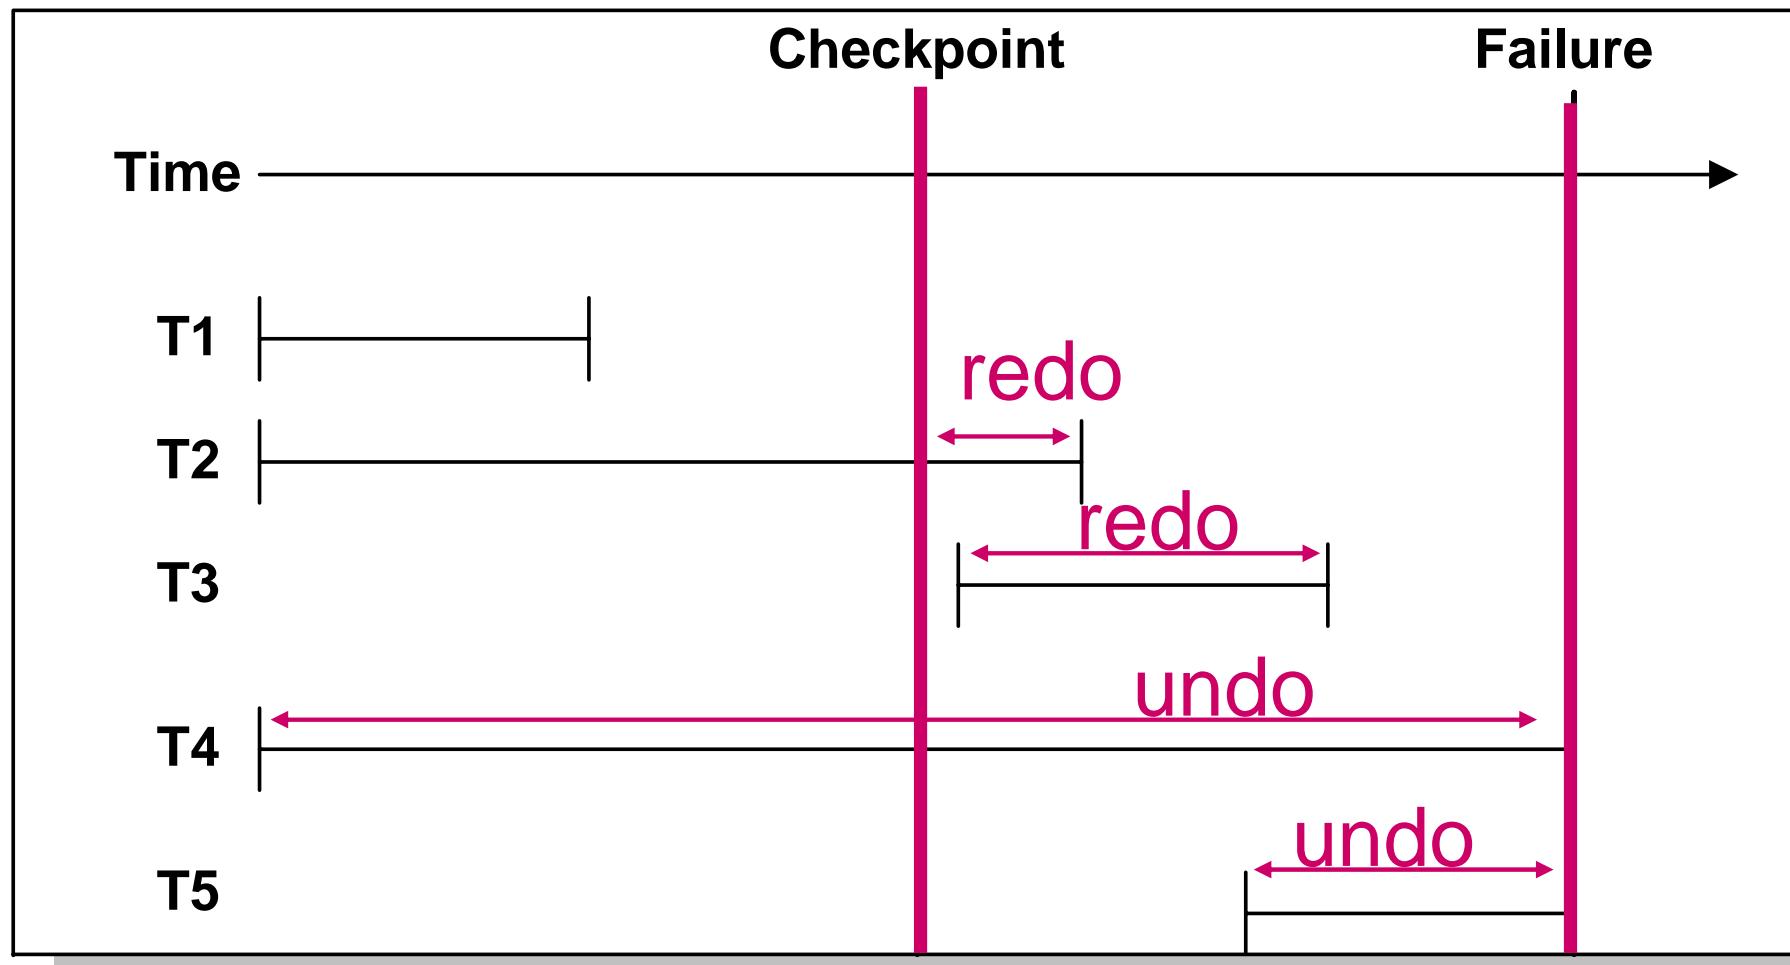
✿ Undo process

✿ send failure message to the computer that running the transaction

✿ rollback

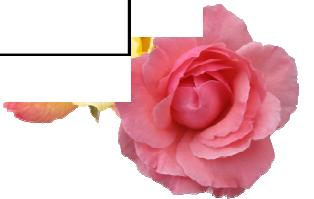


Immediate Update Recovery

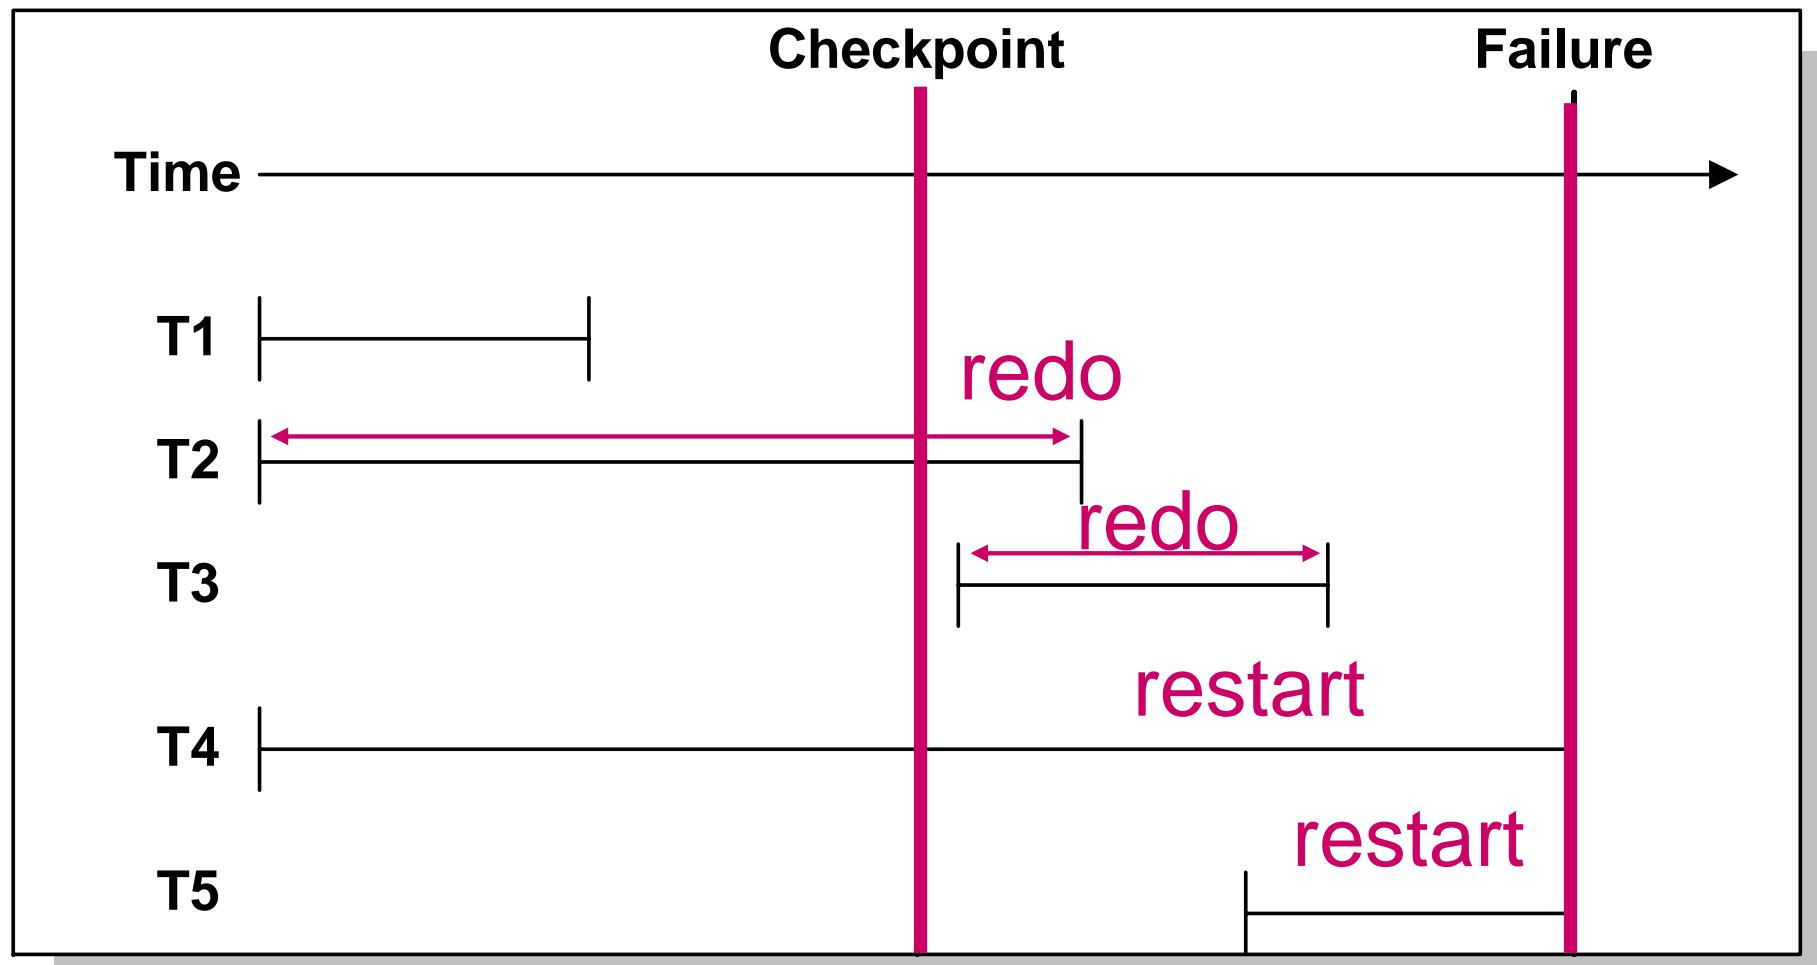


Immediate Update Recovery

Class	Description	Restart Work
T 1	Finished before CP	None
T 2	Started before CP; finished before failure	Redo forward from checkpoint
T 3	Started after CP; finished before failure	Redo forward from checkpoint
T 4	Started before CP; not yet finished	Undo backwards from most recent log record
T 5	Started after CP; not yet finished	Undo backwards from most recent log record



Deferred Update Recovery

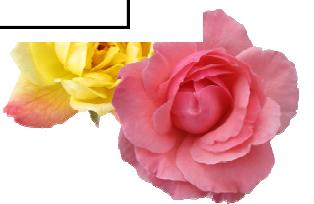


Restart T4 ,T5 without undo



Deferred Update Recovery

Class	Description	Restart Work
T1	Finished before CP	None
T2	Started before CP; finished before failure	Redo forward from first log record
T3	Started after CP; finished before failure	Redo forward from first log record
T4	Started before CP; not yet finished	None
T5	Started after CP; not yet finished	None



Transaction Design Issues

- ✿ Transaction boundary
- ✿ Isolation levels
- ✿ Deferred constraint checking
- ✿ Savepoints



Transaction Boundary Decisions

- ✿ Division of work into transactions
- ✿ Objective: minimize transaction duration
- ✿ Constraint: enforcement of important integrity constraints
- ✿ Transaction boundary decision can affect hot spots



Registration Form Example

Registration Form

Registration No.	1234	Social Security No.	123-45-6789
Status	P	Name	HOMER WELLS
Registration Date	3/29/2005	Address	SEATTLE, WA
Term	Fall	Class	FR
Year	2005		

Enrollments

	Offer No.	Course No.	Units	term	year	location	Time	Instructor
▶	1234	IS320	4	FALL	2005	BLM302	10:30 AM	LEONARD VINCE
*								

Record: * of 1

Total Units Price per Hour
Fixed Charge Total Cost

Record: * of 21

Transaction Boundary Choices

- ✿ One transaction for the entire form
- ✿ One transaction for the main form and one transaction for all subform records
- ✿ One transaction for the main form and separate transactions for each subform record



Avoiding User Interaction Time

- ✿ Avoid to increase throughput
- ✿ Possible side effects: user confusion due to database changes
- ✿ Balance increase in throughput with occurrences of side effects
- ✿ Most situations increase in throughput more important than possible user confusion



Isolation Levels

- ❖ Degree to which a transaction is separated from the actions of other transactions
- ❖ Balance concurrency control overhead with interference problems
- ❖ Some transactions can tolerate uncommitted dependency and inconsistent retrieval problems
- ❖ Specify using the SET TRANSACTION statement



Uncommitted Dependency

✿ or Dirty Read / Temporary Update

✿ occurs when a transaction **is allowed to** retrieve or (worse) update a record that has been updated by another transaction but *has not yet been committed*



Uncommitted Dependency

Time	T ₁	T ₂	Comment
1	read_item(X);		
2	X:=X-N;		
3	write_item(X);		X is temporarily updated
4		read_item(X);	
5		X:=X+M	
6		write_item(X);	
7	read_item(Y);		
...		...	
	ROLLBACK		T ₁ fails and must change the value of X back to its old value; meanwhile T ₂ has read the temporary incorrect value of X

T₂ read x
updated by T₁



Time	T ₁	T ₂	Comment
1	read_item(X);		
2	X:=X-N;		
3	write_item(X);		X is temporarily updated
4		read_item(X);	
5		X:=X+M	
6		write_item(X);	Losted update
7		COMMIT	T ₂ loses update chance
8	read_item(Y);		
...		...	
	ROLLBACK		T ₂ depends on uncommitted value and loses an update at time step 7

it also loses an update at time step 7, because the ROLLBACK in T1 causes data item X to be restored to its value before time step 1.



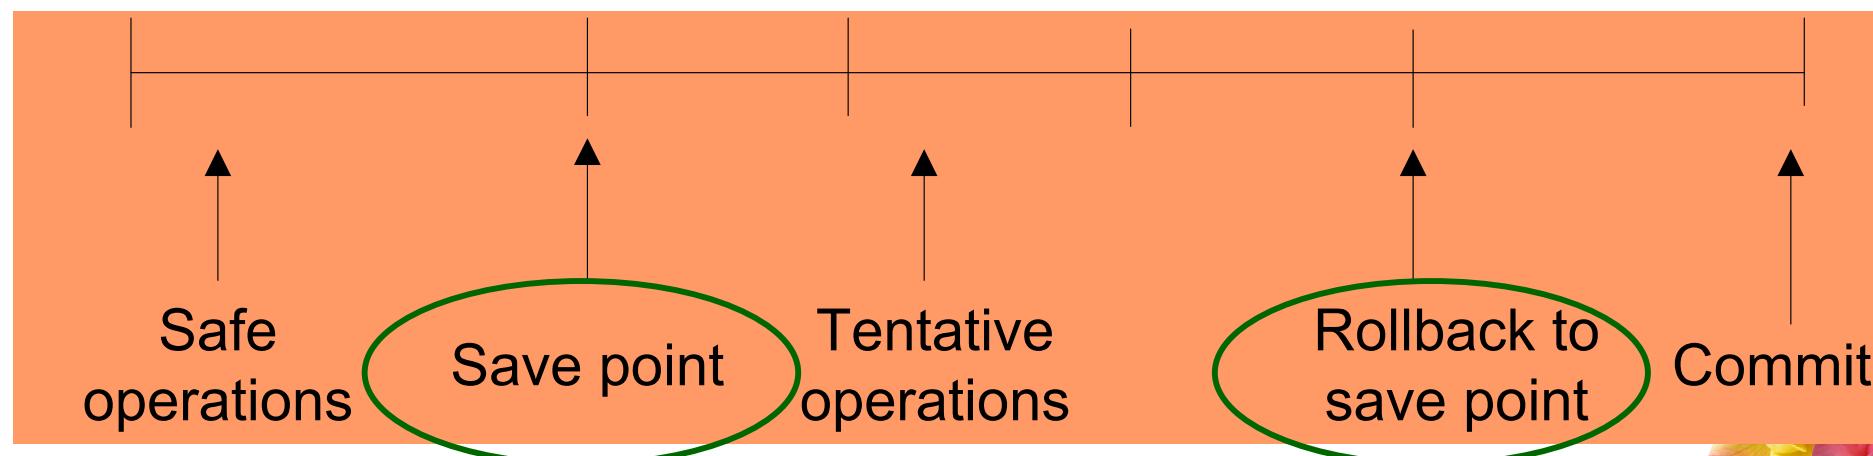
Integrity Constraint Timing

- ✿ Most constraints checked immediately
- ✿ Can defer constraint checking to EOT
- ✿ SQL
 - ✿ Constraint timing clause for constraints in a CREATE TABLE statement
 - ✿ SET CONSTRAINTS statement



Save Points

- ✿ Some transactions have tentative actions
- ✿ **SAVEPOINT** statement determines intermediate points
- ✿ **ROLLBACK** to specified save points



Summary

- ✿ Transaction: user-defined collection of work
- ✿ DBMSs support ACID properties
- ✿ Knowledge of concurrency control and recovery important for managing databases
- ✿ Transaction design issues are important
- ✿ Transaction processing is an important part of workflow management

